Adaptive Sensing beyond Non-Adaptive Information Bounds: End-to-End Co-Design of Geometry, Policy, and Inference via Differentiable Dynamic Programming

Arvin Keshvari, William Tuxbury, Zin Lin
Bradley Department of Electrical and Computer Engineering, Virginia Tech, Blacksburg, VA 24060

April 24, 2026

Abstract

Inverse design has made vast physical parameter spaces a substrate for emergent behavior. In sensing, the stakes of this principle are sharpest at the analog-to-digital boundary, where any information the hardware fails to capture is information no downstream algorithm can recover; hardware optimization alone is therefore not enough, and the geometry must be co-designed with a rule for what to measure next. We formulate this co-design as *joint dynamic programming* (joint-DP): a single optimization over the continuous hardware geometry and a Bellman-optimal adaptive measurement policy. The outer hardware gradient is computed by differentiable dynamic programming with a sharp Bellman maximum, which the envelope theorem makes exact and bias-free, and a relaxation hierarchy carries the common framework from small discrete POMDPs to 10^5-pixel photonic topologies. Across three case studies, joint-DP beats the natural baseline of its community by a large factor: on a radar beam-search POMDP, classical information-bound-guided geometry selection loses 2.8× in attainable adaptive value; on a superconducting-qubit flux sensor, joint-DP reduces deployed mean-squared error by 11.3× over the joint Bayesian Cramér–Rao baseline, with both geometries numerically co-optimized under their respective objectives; and on a 90,000-pixel photonic metasensor, joint-DP via the Bayesian Fisher-information-matrix surrogate reduces deployed mean-squared error by 123× relative to a non-co-designed baseline. For any sensor whose hardware is designed once but whose policy runs for the device's lifetime, joint optimization of hardware and policy is the minimum principled procedure.

Contents

1 Introduction **2**
1.1 Contributions 4
1.2 Relation to prior work 4
1.3 Organization 5

2 Problem setup and the scalar-reward family **5**
2.1 Sensor model 5
2.2 Joint optimization statement 6
2.3 The scalar-reward family Φ 7
2.4 Comparison and operational guidance 8

3 Value decomposition and the ignorance gap **8**
3.1 The four functionals and their ordering 8
3.2 Why joint hardware–policy co-design is necessary 9
3.3 Why EVPI is not a shortcut 9

4 Differentiable dynamic programming for the outer gradient **10**
4.1 Differentiable DP in context . . . 10
4.2 The envelope identity . . . 10
4.3 Unfolding the envelope gradient through the Bellman recursion . . . 11
4.4 Consequences and the outer policy-iteration algorithm . . . 12
4.5 A note on continuous actions . . . 13

5 A hierarchy of relaxations **13**

6 Case Study A: Radar beam search (exact DP, discrete state, $\Phi_{\mathbf{MI}}$) **16**
6.1 Problem setup . . . 16
6.2 Enumeration results . . . 17
6.3 Sweep over the top-hat beam width . . . 17
6.4 Three-way comparison: EVPI picks the wrong geometry . . . 18
6.5 Optimal adaptive policy structure . . . 18

7 Case Study B: Superconducting-qubit flux sensor (exact DP, continuous state, $\Phi_{\mathbf{Var}}$) **18**
7.1 Problem setup . . . 18
7.2 Exact Bellman DP via count-statistic memoization . . . 19
7.3 Joint-DP co-design . . . 19
7.4 Joint PCRB baseline . . . 19
7.5 Deployed MSE comparison . . . 20
7.6 Why $\phi_{\max} = 0.1$: aliasing, horizon, and PCRB validity . . . 20

8 Case Study C: Photonic metasensor (Gaussian-trace surrogate, 10^5-dim continuous $\boldsymbol{c}$, $\Phi_{\log J}$) **20**
8.1 Problem setup . . . 20
8.2 Physics and differentiability . . . 21
8.3 Joint optimization . . . 22
8.4 Optimized binary design and loss trajectory . . . 22
8.5 Monte-Carlo EKF deployment . . . 22
8.6 Feasibility discussion . . . 23

9 Discussion **23**
9.1 The common shape of the adaptive advantage . . . 23
9.2 Why the PCRB baseline loses even when it co-optimizes $\boldsymbol{c}$ . . . 23
9.3 What the EVPI *is* still good for . . . 24
9.4 Limitations . . . 24
9.5 Open directions: dynamical state . . . 24

10 Conclusion **25**

1 Introduction

Modern deep learning has revealed that emergent intelligent behavior can be extracted from optimization over vast parameter spaces: neural networks with billions of weights, trained by gradient descent against differentiable objectives, produce capabilities no hand-designed architecture reaches. A structurally analogous revolution has been unfolding in physics. Photonic crystals engineer electromagnetic band structure through lattice geometries no monatomic medium provides; metasurfaces compress wavefront-shaping into subwavelength films through

the orchestration of millions of sub-wavelength scatterers; inverse-designed silicon photonics [1–6] achieves mode converters, spectral filters, and freeform functional elements in footprints no hand-designed waveguide can reach; acoustic and elastic metamaterials confer effective negative moduli through sub-unit-cell resonances; topological photonics and iridescent biological structures exploit long-range geometric order to produce behavior no constituent chemistry can. The template is the same as deep learning's: a vast parameter space, namely the geometry of a physical medium, optimized by gradient descent against a differentiable physical objective (scattering amplitudes, information content, coherent-addition structure) produces functional behavior no simple structure can. In this sense, *structural complexity in physics is the analog of representational complexity in deep networks*: both are substrates through which optimization converts raw degrees of freedom into engineered capability. The consequence is that physical hardware optimization is not a peripheral engineering concern but a fundamental question about where intelligence can live.

This question sharpens dramatically in sensing and metrology, at the analog-to-digital boundary where a physical signal becomes a digital record and any information lost there bounds everything any downstream algorithm can recover. The conventional blueprint of the last two decades pairs a generic, task-agnostic transducer front-end with powerful digital reconstruction, and so treats the physical hardware as a passive conduit and asks the digital stages to rescue whatever structure the geometry happens to impose. Computational inverse design already inverts this blueprint on the hardware side, routinely saturating the electromagnetic performance bounds of the underlying scattering problem [3, 5, 7, 8], and end-to-end optical co-design extends it by tying the geometry to a downstream reconstruction network, closing the loop from physics to inference in a single gradient [9–18]. What remains missing is the policy dimension: the recognition that the geometry must co-design with *a rule for what to measure next*, i.e., a multi-step adaptive measurement policy that reacts to intermediate data. When such a policy is installed, part of the intelligence-extraction work migrates upstream into the analog domain before digitization, and the co-optimized geometry becomes a participant in the inference itself rather than a conduit to it. This paper supplies that missing dimension.

A mature literature on adaptive sensing and sequential Bayesian experimental design constructs formally-optimal multi-step measurement policies for partially-observable Markov decision processes (POMDPs) [19–31], but treats the physical hardware as immutable. The hardware half (inverse design, end-to-end optics) and the policy half (adaptive-sensing POMDP theory) are thus designed separately, each community optimizing its own half against the other's default: a textbook setup for a joint optimum that nobody reaches. Several near-relatives live in the intersection without closing it: cognitive radar adapts the waveform on a fixed antenna [32–35], Kitaev-style quantum phase estimation adapts the pulse sequence on a fixed circuit [36–42], programmable-metasurface imaging adapts the reconstruction on a fixed scattering medium [43, 44], and variational quantum sensing parametrizes the probe through a trainable neural ansatz rather than a geometric topology [45, 46]. None of these treats continuous physical geometry and a multi-step Bellman-optimal measurement policy as co-equal design variables in a single optimization.

Closing this gap requires a framework that (a) is cleanly defined across discrete and continuous state spaces, (b) admits a tractable outer gradient with respect to the physical geometry, and (c) offers a hierarchy of relaxations scaling from exact dynamic programming on small POMDPs to high-dimensional continuous hardware inverse design. We supply all three, and call the procedure *joint dynamic programming*, or *joint-DP*: a single optimization that treats the continuous hardware $\boldsymbol{c}$ and the Bellman-optimal adaptive policy π as co-equal decision variables, with π re-solved by dynamic programming at each outer iterate in $\boldsymbol{c}$. To our knowledge, this is the first formulation in which continuous physical hardware and a multi-step Bellman-optimal adaptive policy are co-designed as a single optimization problem.

Joint-DP improves on the two natural baselines that bracket the problem, on benchmarks

where exact comparison is tractable. The adaptive-sensing POMDP literature runs the Bellman dynamic-programming (DP) recursion — the standard backward-in-time value-iteration procedure that solves for an optimal measurement policy at fixed hardware. Joint-DP strictly generalizes this by co-optimizing $\boldsymbol{c}$ under the adaptive policy, yielding a 123× deployed-MSE improvement on the photonic metasensor of Case Study C over a non-co-designed baseline. The OED and physics-of-bounds literatures select $\boldsymbol{c}$ to maximize a prior-averaged, fixed-schedule information functional (mutual information, expected-Fisher log-determinant, Expected Value of Perfect Information) and attach an adaptive policy afterwards; on the radar POMDP of Case Study A the information-functional-maximizing $\boldsymbol{c}$ loses a factor of 2.8× in attainable adaptive value to the joint-DP optimum, and on the superconducting-qubit flux sensor of Case Study B the joint Bayesian Cramér–Rao fixed-schedule baseline loses a factor of 11.3× in deployed Bayesian MSE ($z = +132\sigma$). The performance gaps are first-order in engineering terms.

1.1 Contributions

1. **Unified formulation.** A family of pointwise scalar rewards $\Phi(\boldsymbol{x}, \boldsymbol{s}, \boldsymbol{c})$ (mutual-information-aligned Φ_{MI} [27], Bayesian-MSE-aligned Φ_{Var}, and Fisher-log-det $\Phi_{\log J}$ [47, 48]) fit the same joint-$(\boldsymbol{c}, \pi)$ Bellman machinery. The pointwise dependence on the unknown state $\boldsymbol{x}$ is the native reward form of POMDP theory, but it is rarely foregrounded in the adaptive-sensing and Bayesian-OED communities, which typically work directly with prior-averaged belief functionals; foregrounding it is what unifies three traditionally separate reward families under one umbrella and what makes the oracle / fixed / adaptive decomposition of §3 well-defined. Choosing Φ is a modeling commitment set by the problem's deployment metric.

2. **A sharp-max, bias-free differentiable DP for belief-state POMDPs.** We compute the outer hardware gradient of a POMDP Bellman recursion by keeping the max operator sharp throughout and appealing to the envelope theorem [49, 50] at the Bellman-optimal argmax; reverse-mode automatic differentiation through an argmax-frozen version of the recursion returns an exact, unbiased hardware gradient with no smoothing temperature. This differs from the dominant softmax-regularized and implicit-KKT lines of differentiable DP [51–54] in both method (no smoothing of the max; no re-formulation of the inner problem via KKT surrogates) and setting (belief-state POMDPs with Bayesian measurement dynamics, rather than fully-observed MDPs or structured-prediction trellises); to our knowledge the combination is new to the differentiable-DP literature.

3. **Quantitative findings.** On benchmarks where exact comparison is tractable, joint-DP delivers performance gaps of 2.8× against classical information-bound-guided geometry selection (radar, Case Study A); 11.3× at $z = +132\sigma$ in deployed Bayesian MSE against the joint Bayesian Cramér–Rao baseline [55, 56] (scqubit, Case Study B); and 123× in deployed MSE against a non-co-designed baseline (photonic, Case Study C).

4. **Portability across scales.** A relaxation cascade (exact DP, certainty-equivalent, myopic one-step, Fisher-trace plus EKF) carries the framework from small POMDPs to a 90,000-pixel photonic metasensor, with each rung's failure mode named.

1.2 Relation to prior work

The components of our framework are individually well-established: exact Bellman DP for POMDPs [20, 21, 57]; mutual information as an experimental-design reward (Lindley [27]); end-to-end optical co-design with learned reconstruction, including information-reward end-to-end imaging by Waller and collaborators [58–61]; differentiable dynamic programming as a gradient-propagation vehicle [51–54]; the envelope/Danskin theorem [49, 50, 62]; Expected Value of

Perfect Information as a classical decision-theoretic quantity [63, 64]. The joint optimization of *continuous* physical hardware with a *multi-step Bellman-optimal* adaptive policy is not standard in any of these: end-to-end optics uses one-step reconstructions, adaptive sensing fixes the hardware, and the EVPI/IG line does not treat continuous physical geometry as a first-class design variable. Each ingredient is classical; the novelty is their synthesis into a single joint optimization over continuous physical hardware and a multi-step Bellman-optimal adaptive policy, with sharp-max differentiable DP as the vehicle for the outer gradient.

Our framework sits downstream of, and is complementary to, the recent program on fundamental limits to light–matter interaction via convex-duality and T-matrix methods [7, 8, 65–76]: those bounds certify whether a target is achievable by *any* passive medium; we select the $\boldsymbol{c}$ that realizes it along with the policy that consumes it. Likewise for the quantum Fisher information program on wave and photon scattering [77–83]: our $\Phi_{\log J}$ choice in the Gaussian-posterior regime is the Bayesian analog of the maximum-Fisher-information-state object; outside that regime it shares the stage with alternative scalar rewards chosen by the problem's deployment metric.

1.3 Organization

The paper builds up from formulation to evidence. Section 2 states the joint-$(\boldsymbol{c}, \pi)$ optimization and develops the family of pointwise scalar rewards Φ. Section 3 introduces the value decomposition and the ignorance gap. Section 4 develops differentiable dynamic programming as the vehicle for the outer gradient; Section 5 catalogs the relaxation hierarchy. Sections 6, 7, and 8 are the three case studies (discrete POMDP; continuous-state qubit sensor; high-dimensional photonic metasensor). Sections 9–10 discuss limitations and conclude.

2 Problem setup and the scalar-reward family

2.1 Sensor model

A *sensor* in this paper is a tuple $(\boldsymbol{c}, \pi)$ where $\boldsymbol{c} \in \mathcal{C} \subseteq \mathbb{R}^{d_c}$ is a continuous *physical geometry* fixed at manufacturing time (the flux-bias and readout-line mutual inductances of a SQUID loop, the permittivity distribution of an inverse-designed metasurface, the elements of a phased-array antenna), and $\pi : \mathcal{B} \to \mathcal{S}$ is a *measurement policy* that maps the current belief $\boldsymbol{b} \in \mathcal{B}$ over the unknown state $\boldsymbol{x} \in \mathcal{X}$ to an action $\boldsymbol{s} \in \mathcal{S}$. At each step $k \in \{1, \ldots, K\}$ the policy emits an action $\boldsymbol{s}_k = \pi(\boldsymbol{b}_k)$, whose execution produces a *measurement* $\boldsymbol{y}_k \in \mathcal{Y}$ drawn from the likelihood

$$\boldsymbol{y}_k \;\sim\; p(\boldsymbol{y}_k \mid \boldsymbol{x}, \boldsymbol{s}_k;\, \boldsymbol{c}), \tag{1}$$

where $\mathcal{Y}$ is the measurement space specified by the physical readout (discrete detector counts in a photon-starved regime, photon clicks in a Ramsey sequence, real-valued voltages from a homodyne detector, etc.) and $\boldsymbol{c}$ parametrizes the likelihood through its physical forward model (Maxwell's equations, circuit Hamiltonian dynamics, Schrödinger evolution with pulse shaping, etc.). For non-adaptive ("fixed-schedule") policies π is a constant function of the belief; for adaptive policies π depends on the history $\mathcal{F}_{k-1} = \{\boldsymbol{s}_{1:k-1}, \boldsymbol{y}_{1:k-1}\}$ through the posterior $\boldsymbol{b}_k = p(\boldsymbol{x} \mid \mathcal{F}_{k-1};\, \boldsymbol{c})$. We write K for the *horizon* — the total number of measurement steps the operator is budgeted before an estimate of $\boldsymbol{x}$ is reported, which may be limited by coherence time in a quantum sensor, by integration-time budget in a radar search, or by throughput requirements in an imaging pipeline — $p_0(\boldsymbol{x})$ for the prior, and reserve $\pi^\star(\boldsymbol{c})$ for the Bellman-optimal policy at fixed $\boldsymbol{c}$. Inner expectations are over observations $\boldsymbol{y}$ with $\boldsymbol{x}, \boldsymbol{c}$ held fixed; outer expectations are over the prior p_0.

2.2 Joint optimization statement

A *joint co-design objective* is any scalar functional of the hardware-policy pair,

$$V(\boldsymbol{c}, \pi) \;=\; \mathbb{E}_{\boldsymbol{x}\sim p_0}\, \mathbb{E}_{\boldsymbol{y}_{1:K}|\boldsymbol{x},\pi,\boldsymbol{c}}\big[\, \Phi(\boldsymbol{x}, \boldsymbol{s}^{\pi}_{1:K}, \boldsymbol{c})\,\big], \tag{2}$$

where $\Phi : \mathcal{X} \times \mathcal{S}^K \times \mathcal{C} \to \mathbb{R}$ is a user-specified *pointwise reward*: "pointwise" in the sense that Φ is evaluated at each realization of the true state $\boldsymbol{x}$, action sequence $\boldsymbol{s}_{1:K}$, and hardware $\boldsymbol{c}$ separately, before any averaging over the prior. Aggregate objectives (Bayesian mean-squared error, mutual information, log-determinant of an expected Fisher matrix) are recovered by taking the expectations in Equation (2). We correspondingly use "pointwise MSE" for the state-conditional mean-squared error $\mathbb{E}_{\boldsymbol{y}|\boldsymbol{x}}[(\hat{\boldsymbol{x}} - \boldsymbol{x})^2]$, which is an $\boldsymbol{x}$-dependent scalar, as opposed to the aggregate Bayesian MSE which is its prior average. The joint hardware–policy optimization solves

$$(\boldsymbol{c}^\star, \pi^\star) \;=\; \underset{\boldsymbol{c}\in\mathcal{C},\ \pi\in\Pi}{\arg\max}\; V(\boldsymbol{c}, \pi). \tag{3}$$

Equation (3) is the sole formal statement in the paper; what varies across applications is only the choice of Φ. Two degenerate cases recover classical programs: $\Pi = \{\text{constant policies}\}$ with Φ minus a reconstruction loss recovers **end-to-end optics** [9, 10]; $\mathcal{C} = \{\boldsymbol{c}_0\}$ (fixed hardware) with any Φ recovers **adaptive sensor management** [19, 22]. Their pairing is the content of this paper.

For the posterior-mean estimator on continuous $\boldsymbol{x}$, the tower-law identity $\mathbb{E}[(\hat{\boldsymbol{x}}_{\mathrm{pm}} - \boldsymbol{x})^2] = \mathbb{E}[\mathrm{Var}(\boldsymbol{x} \mid \boldsymbol{b}_K)]$ makes the Bayesian MSE a special case of $V(\boldsymbol{c}, \pi)$ with $\Phi = -\mathrm{Var}(\boldsymbol{x} \mid \boldsymbol{b}_K)$. Section 2.3 develops this and two alternative Φ choices used by the other case studies.

Two-level structure (joint-DP). When Π is the full space of adaptive policies, the inner $\arg\max_\pi$ in Equation (3) is solved by the Bellman recursion on the belief state $\boldsymbol{b}_k$ (full derivation from Equation (2) in the accompanying supplement),

$$V_K(\boldsymbol{b};\, \boldsymbol{c}) \;=\; \mathbb{E}_{\boldsymbol{x}|\boldsymbol{b}}[\, \Phi_0(\boldsymbol{x}, \boldsymbol{b};\, \boldsymbol{c})\,], \tag{4}$$

$$V_k(\boldsymbol{b};\, \boldsymbol{c}) \;=\; \max_{\boldsymbol{s}_k\in\mathcal{S}}\; \mathbb{E}_{\boldsymbol{y}_k|\boldsymbol{b},\boldsymbol{s}_k,\boldsymbol{c}}\big[\, V_{k+1}\big(T(\boldsymbol{b}, \boldsymbol{s}_k, \boldsymbol{y}_k;\, \boldsymbol{c});\, \boldsymbol{c}\big)\,\big], \quad k = K-1, \ldots, 0, \tag{5}$$

where $T(\boldsymbol{b}, \boldsymbol{s}, \boldsymbol{y};\, \boldsymbol{c})$ is the Bayesian belief-update operator,

$$T(\boldsymbol{b}, \boldsymbol{s}, \boldsymbol{y};\, \boldsymbol{c})(\boldsymbol{x}) \;=\; \frac{p(\boldsymbol{y} \mid \boldsymbol{x}, \boldsymbol{s};\, \boldsymbol{c})\, \boldsymbol{b}(\boldsymbol{x})}{\int_{\mathcal{X}} p(\boldsymbol{y} \mid \boldsymbol{x}', \boldsymbol{s};\, \boldsymbol{c})\, \boldsymbol{b}(\boldsymbol{x}')\, d\boldsymbol{x}'}, \tag{6}$$

the posterior over $\boldsymbol{x}$ obtained from prior $\boldsymbol{b}$ by Bayes' rule given one measurement $\boldsymbol{y}$ taken at action $\boldsymbol{s}$ under hardware $\boldsymbol{c}$ (the integral in the denominator is a sum for discrete $\boldsymbol{x}$). $\Phi_0(\boldsymbol{x}, \boldsymbol{b}_K;\, \boldsymbol{c})$ is the terminal pointwise reward written as a function of the terminal belief; for $\Phi_{\log J}$ the belief is augmented by the accumulated Fisher-information matrix, and Φ_0 is read off the augmentation. The inner optimal policy at each step is $\pi^\star_k(\boldsymbol{b};\, \boldsymbol{c}) = \arg\max_{\boldsymbol{s}_k} \mathbb{E}_{\boldsymbol{y}_k}[V_{k+1}(T(\boldsymbol{b}, \boldsymbol{s}_k, \boldsymbol{y}_k;\, \boldsymbol{c});\, \boldsymbol{c})]$, and the joint hardware-policy optimum is the outer maximum of the inner DP value at the initial belief $\boldsymbol{b}_0 = p_0$,

$$\boxed{\boldsymbol{c}^\star \;=\; \underset{\boldsymbol{c}\in\mathcal{C}}{\arg\max}\; V^\star(\boldsymbol{c}), \qquad V^\star(\boldsymbol{c}) \;\coloneqq\; V_0(\boldsymbol{b}_0;\, \boldsymbol{c}).} \tag{7}$$

We call Equations (4)–(7) the *joint-DP* structure: an outer continuous optimization over the physical geometry $\boldsymbol{c}$ wrapped around an inner Bellman recursion that produces the optimal adaptive policy $\pi^\star(\boldsymbol{c})$ at each candidate $\boldsymbol{c}$. Section 4 develops differentiable dynamic programming as the vehicle for the outer $\boldsymbol{c}$-gradient, which decouples from the inner Bellman solve via the envelope theorem; Section 5 describes a cascade of relaxations for the inner recursion when exact DP is infeasible.

2.3 The scalar-reward family Φ

Three canonical choices for Φ emerge naturally from three different bodies of theory. In each case Φ is a deterministic function of $(\boldsymbol{x}, \boldsymbol{s}_{1:K}, \boldsymbol{c})$; any expectation over the measurement trajectory $\boldsymbol{y}_{1:K}$ that appears in the definition is baked in, so that the outer objective $V(\boldsymbol{c}, \pi) = \mathbb{E}_{\boldsymbol{x},\boldsymbol{y}}[\Phi]$ of Equation (2) is a single expectation over the prior and the data.

Mutual-information-aligned reward Φ_{MI}. Following the Shannon–Lindley line [27, 84] through Chaloner–Verdinelli [31], Sebastiani–Wynn [85], and the adaptive sensor-management literature [22, 23],

$$\Phi_{\mathrm{MI}}(\boldsymbol{x}, \boldsymbol{s}, \boldsymbol{c}) \;=\; \mathbb{E}_{\boldsymbol{y}|\boldsymbol{x},\boldsymbol{s},\boldsymbol{c}}\big[\ln|\mathrm{supp}(\boldsymbol{b}_K)| - H(\boldsymbol{b}_K)\big], \tag{8}$$

the expected negative entropy of the terminal belief $\boldsymbol{b}_K$ at state $\boldsymbol{x}$. Averaging over the prior gives $\mathbb{E}_{\boldsymbol{x}}[\Phi_{\mathrm{MI}}] = I(\boldsymbol{x}; \boldsymbol{y}_{1:K} \mid \boldsymbol{s}, \boldsymbol{c})$, the mutual information. This is the reward of choice for discrete-state problems where MSE is undefined or dominated by 0/1 classification loss [24], and for any problem in the large-K Bernstein–von Mises regime [86, 87] (the asymptotic limit in which the posterior converges to a Gaussian centered near the true parameter, making MI monotone in MSE). The photonic and quantum-sensing communities have developed extensive machinery for *upper bounds* on achievable mutual information given physical constraints (reciprocity, passivity, energy); our framework consumes those bounds as constraints on the feasible set $\{\Phi_{\mathrm{MI}}(\cdot, \cdot, \boldsymbol{c}) : \boldsymbol{c} \in \mathcal{C}\}$ rather than re-deriving them.

Bayesian-MSE-aligned reward Φ_{Var}. Inheriting from the Bayesian Cramér–Rao tradition [55] and the posterior-variance-minimization line of Bayesian optimal experimental design [31, 88], the tower law of variance (the identity that the Bayes-averaged squared error of the posterior mean equals the Bayes-averaged posterior variance)

$$\mathbb{E}_{\boldsymbol{x},\boldsymbol{y}_{1:K}}\big[(\hat{\boldsymbol{x}}_{\mathrm{pm}} - \boldsymbol{x})^2\big] \;=\; \mathbb{E}_{\boldsymbol{y}_{1:K}}[\mathrm{Var}(\boldsymbol{x} \mid \boldsymbol{y}_{1:K})] \tag{9}$$

(with $\hat{\boldsymbol{x}}_{\mathrm{pm}}$ the posterior mean) motivates

$$\Phi_{\mathrm{Var}}(\boldsymbol{x}, \boldsymbol{s}, \boldsymbol{c}) \;=\; \mathbb{E}_{\boldsymbol{y}|\boldsymbol{x},\boldsymbol{s},\boldsymbol{c}}\big[-\mathrm{Var}(\boldsymbol{x} \mid \boldsymbol{b}_K)\big], \tag{10}$$

for which the outer objective $V(\boldsymbol{c}, \pi)$ equals $-1\times$ the deployed Bayesian MSE exactly. The joint-DP optimum of this reward is literally the *MSE-optimal* adaptive policy at fixed $\boldsymbol{c}$, without any asymptotic approximation. For continuous-state, finite-K problems with bounded-support priors, where the Bernstein–von Mises regime is not reached, Φ_{Var} strictly dominates Φ_{MI} as an MSE anchor. The tower-law identity is textbook; its use as a joint-DP terminal reward in the present hardware-policy setting is new.

Fisher-information-aligned reward $\Phi_{\log J}$. Inheriting from the classical Fisher-information-matrix tradition [47, 48, 89] and its Bayesian extension [55, 56], and applicable to continuous $\boldsymbol{x}$ with a Gaussian-posterior surrogate,

$$\Phi_{\log J}(\boldsymbol{x}, \boldsymbol{s}, \boldsymbol{c}) \;=\; \tfrac{1}{2}\log\det J_N(\boldsymbol{x}, \boldsymbol{s}_{1:K}, \boldsymbol{c}), \tag{11}$$

where J_N is the accumulated Fisher information matrix at true $\boldsymbol{x}$. The outer expectation is the log-determinant of the Bayesian Fisher information matrix [31, 90], with J_N^{-1} the Bayesian CRB. In the Gaussian-posterior regime $\Phi_{\log J}$ agrees with Φ_{Var} up to log-transforms. Outside that regime (biased estimators, bounded priors, small K), $\Phi_{\log J}$ retains an asymptotic CRB interpretation while losing its direct link to finite-sample MSE. It is the natural reward for high-dimensional continuous problems where the full Bellman DP is intractable and the Gaussian surrogate is the cheapest defensible relaxation. The quantum-Fisher-information program

on wave and photon scattering [77–79, 81] and the photonic information-bound program on scattering-matrix extremization [7, 8, 65–67] establish upper bounds on $\det J_N$ over admissible $\boldsymbol{c}$; they tell us what $\sup_{\boldsymbol{c}} \det J_N$ is, and our framework tells us which $\boldsymbol{c}$ and which policy realize it jointly.

2.4 Comparison and operational guidance

Table 1: Three natural pointwise scalar rewards, their operational semantics, and their alignment with finite-K MSE.

reward	terminal expression	community	tied to pointwise MSE?
Φ_{MI}	$\ln \lvert\mathrm{supp}(\boldsymbol{b}_K)\rvert - H(\boldsymbol{b}_K)$	sensor management, sOED	asymptotic only
Φ_{Var}	$-\mathrm{Var}(\boldsymbol{x} \mid \boldsymbol{b}_K)$	Bayesian estimation	**directly, by tower law**
$\Phi_{\log J}$	$\frac{1}{2}\log\det J_N(\boldsymbol{x}, \boldsymbol{s}, \boldsymbol{c})$	Fisher info, PCRB	asymptotic only

The three rewards agree asymptotically under Bernstein–von Mises but diverge at finite horizons, and the choice of Φ matters operationally for the co-design optimum at modest horizon on a bounded prior. We accordingly present the paper's machinery as reward-agnostic: the user picks Φ based on the problem's state space, horizon, and deployment metric, and the differentiable-DP outer gradient of Section 4 derives identically from any of them.

3 Value decomposition and the ignorance gap

For any choice of Φ, the outer objective $V(\boldsymbol{c}, \pi)$ admits a four-functional decomposition that is useful for both intuition and diagnostics. The decomposition is standard in sequential Bayesian experimental design [28, 29] and was used in the radar sensor-management literature under the name "oracle value" of perfect information [22, 91]; we restate it in the Φ-agnostic form used in the rest of the paper.

3.1 The four functionals and their ordering

$$V_{\mathrm{oracle}}(\boldsymbol{x}, \boldsymbol{c}) = \max_{\boldsymbol{s}_{1:K}} \Phi(\boldsymbol{x}, \boldsymbol{s}_{1:K}, \boldsymbol{c}), \tag{12}$$

$$V_{\mathrm{fixed}}(\boldsymbol{c}) = \max_{\boldsymbol{s}_{1:K}} \mathbb{E}_{\boldsymbol{x}}\big[\Phi(\boldsymbol{x}, \boldsymbol{s}_{1:K}, \boldsymbol{c})\big], \tag{13}$$

$$V_{\mathrm{adaptive}}(\boldsymbol{c}) = \max_{\pi \in \Pi} \mathbb{E}_{\boldsymbol{x}}\mathbb{E}_{\boldsymbol{y}|\boldsymbol{x},\pi,\boldsymbol{c}}\big[\Phi(\boldsymbol{x}, \boldsymbol{s}^{\pi}_{1:K}, \boldsymbol{c})\big], \tag{14}$$

$$\mathbb{E}[\mathrm{IG}](\boldsymbol{c}) = \mathbb{E}_{\boldsymbol{x}}[V_{\mathrm{oracle}}(\boldsymbol{x}, \boldsymbol{c})] - V_{\mathrm{fixed}}(\boldsymbol{c}). \tag{15}$$

$V_{\mathrm{oracle}}(\boldsymbol{x}, \boldsymbol{c})$ is the value attained by a *clairvoyant* sensor that knows the true $\boldsymbol{x}$ and picks the best fixed schedule for it. $V_{\mathrm{fixed}}(\boldsymbol{c})$ is the best non-adaptive schedule, $V_{\mathrm{adaptive}}(\boldsymbol{c})$ is the best adaptive policy (the quantity the user ultimately deploys), and $\mathbb{E}[\mathrm{IG}](\boldsymbol{c})$ — the *expected ignorance gap* or *expected value of perfect information* (EVPI) [63, 64] — measures how much adaptive slack exists at that geometry.

The three values are ordered at every $\boldsymbol{c}$,

$$V_{\mathrm{fixed}}(\boldsymbol{c}) \;\le\; V_{\mathrm{adaptive}}(\boldsymbol{c}) \;\le\; \mathbb{E}_{\boldsymbol{x}}[V_{\mathrm{oracle}}(\boldsymbol{x}, \boldsymbol{c})], \tag{16}$$

because fixed schedules are a subset of adaptive policies (left) and no policy can beat the clairvoyant oracle (right). Subtracting V_{fixed} from Equation (16) bounds the *adaptive gain* $V_{\mathrm{adaptive}}(\boldsymbol{c}) - V_{\mathrm{fixed}}(\boldsymbol{c})$ above by the EVPI:

$$0 \;\le\; V_{\mathrm{adaptive}}(\boldsymbol{c}) - V_{\mathrm{fixed}}(\boldsymbol{c}) \;\le\; \mathbb{E}[\mathrm{IG}](\boldsymbol{c}). \tag{17}$$

If this upper bound were tight, the EVPI-maximizing geometry would coincide with the adaptive-gain-maximizing one — the hope that motivates the EVPI as a candidate design criterion. §3.3 shows the hope does not survive scrutiny.

3.2 Why joint hardware–policy co-design is necessary

The quantity the user deploys is $V_{\text{adaptive}}(\boldsymbol{c})$, so the question is which hardware maximizes that deployed value. Two facts make it a function of $\boldsymbol{c}$ that cannot be maximized piecewise.

Fact 1: adaptation extracts real value at most geometries. A fixed schedule commits all K measurements in advance, before any data has arrived. An adaptive policy sees each outcome as it happens and re-routes subsequent measurements accordingly: effort that would have been wasted confirming something the data has already settled gets redirected toward parts of the state space that remain ambiguous. In an imaging problem this is the difference between a fixed raster scan and one that zooms in on the suspected region; in a qubit-flux sensor it is the difference between a fixed Ramsey delay ladder and one that shortens τ when the posterior is broad (to avoid phase-wrapping ambiguity) and lengthens τ once the flux is roughly localized. The adaptive gain of Equation (17) is therefore typically strictly positive and often sizable.

Fact 2: the adaptive gain depends nontrivially on $\boldsymbol{c}$. A hardware configuration that maximizes V_{fixed} is not in general the hardware that maximizes V_{adaptive}: a geometry good for blind schedules is good because its expected measurements are individually informative on average, and that property tends to leave little room for outcome-dependent re-routing. Conversely, a geometry with very large oracle value (the right-hand side of Equation (16)) may put so much of that value out of reach of any policy that only sees past data that the realizable adaptive value is actually small. The geometry that deploys well under an adaptive policy is an *interior* geometry, one that neither is over-tuned to blind schedules nor relies on clairvoyance; this optimum is not located by examining V_{fixed} or V_{oracle} separately.

Consequence: the two layers cannot be solved in sequence. One cannot pick $\boldsymbol{c}$ first (by any single-functional criterion) and then pick π, because the criterion for a good $\boldsymbol{c}$ depends on which π will follow. Joint optimization of $(\boldsymbol{c}, \pi)$ is necessary; the rest of the paper is about making it tractable. Section 4 shows how differentiable dynamic programming (with the envelope theorem providing the sharp-max exactness certificate) decouples the *gradient computation* from the inner-policy solve even when the two quantities themselves do not decouple, so that an outer-loop optimizer can climb $V_{\text{adaptive}}(\boldsymbol{c})$ with a single Bellman solve per iteration.

3.3 Why EVPI is not a shortcut

It is worth asking whether the EVPI can stand in for V_{adaptive} as a design criterion. Since V_{oracle} and V_{fixed} are computable by direct schedule enumeration (no Bellman recursion, no belief tree), $\mathbb{E}[\text{IG}](\boldsymbol{c})$ is a cheap scalar functional of $\boldsymbol{c}$, unlike $V_{\text{adaptive}}(\boldsymbol{c})$ which requires the full backward-in-time value iteration.

The shortcut fails: the upper bound in Equation (17) is slack, and its slackness pattern is non-monotone in $\boldsymbol{c}$, so

$$\arg\max_{\boldsymbol{c}} \mathbb{E}[\text{IG}](\boldsymbol{c}) \neq \arg\max_{\boldsymbol{c}} V_{\text{adaptive}}(\boldsymbol{c}). \tag{18}$$

Extreme geometries (a very narrow beam, a near-isotropic beam) produce high EVPI because the oracle's advantage over any blind fixed schedule is enormous, but the adaptive policy — which does not know $\boldsymbol{x}$ — cannot replicate the oracle's strategy and captures only a small

fraction of the nominal headroom. Interior geometries have a modest EVPI ceiling but let the adaptive policy capture a larger fraction of it, which is what ultimately matters. On the radar benchmark of Section 6, the EVPI-maximizing geometry loses a factor of 2.8× in attainable V_{adaptive} to the joint-DP optimum; the joint solve is unavoidable.

EVPI retains useful roles as a workflow tool rather than a design target: a cheap pre-screen when schedule enumeration is tractable but the full joint DP is not; a post-hoc gap-to-oracle diagnostic at $\boldsymbol{c}^\star$; and an interpretive lens on the shape of the $\boldsymbol{c}$-landscape. §9 develops these.

4 Differentiable dynamic programming for the outer gradient

4.1 Differentiable DP in context

The outer joint-DP optimization (Equation (7)) requires a gradient of the inner Bellman value $V^\star(\boldsymbol{c}) = V_0(p_0; \boldsymbol{c})$ with respect to the continuous hardware $\boldsymbol{c}$. The inner Bellman recursion (Equations (4)–(5)) contains a max over actions at every layer, and max is non-smooth. Differentiating a nested max-over-expectations recursion is therefore the load-bearing operation, and it is the subject of a now-active ML literature under the name *differentiable dynamic programming* [51–54]. Two strategies for handling the max dominate the landscape:

Smooth-max (the dominant approach). Replace max with a smooth relaxation — softmax, Tsallis entropy-regularization, Boltzmann averaging — so the recursion becomes everywhere-differentiable and reverse-mode automatic differentiation handles the rest. This is the path taken by Mensch & Blondel [51], differentiable-MPC work [52], and a broader family of implicit-differentiation methods [53, 54]. The price is a *regularization bias*: the smoothed recursion no longer evaluates $V^\star(\boldsymbol{c})$ exactly, and the gradient is biased by an amount controlled by the smoothing temperature. The bias can be driven to zero by annealing the temperature, but the limit requires a schedule that is itself a hyperparameter.

Sharp-max via envelope (the approach of this paper). Keep the max operator sharp in the recursion and use the envelope theorem [49, 50, 62, 92] to certify that, at the Bellman-optimal argmax, the indirect contribution of argmax sensitivity to $\partial V^\star/\partial \boldsymbol{c}$ vanishes. The recursion keeps its exact Bellman value, and reverse-mode AD through an *argmax-frozen* version of the recursion returns an exact, unbiased hardware gradient with no smoothing bias and no temperature schedule.

The two strategies trade smoothness for exactness. Smooth-max gives everywhere-differentiable code and a biased gradient; sharp-max gives exact gradients and requires one additional structural step (freezing the argmax) to take advantage of the envelope identity. For the physical-hardware gradients this paper is concerned with — where biases in $\partial V^\star/\partial \boldsymbol{c}$ could steer the outer optimization to the wrong $\boldsymbol{c}^\star$ and where the inner Bellman DP is already exact on its own terms — the sharp-max path is the natural choice. Two further distinctions from the existing differentiable-DP literature are worth noting: we work on belief-state POMDPs rather than fully-observed MDPs or structured-prediction trellises, and the hardware parameter $\boldsymbol{c}$ enters both the observation likelihood and the Bayesian belief update T, rather than parametrizing transition costs or trellis scores directly. The remainder of this section develops the sharp-max path in that setting: the envelope identity (§4.2), its operational form unfolded through the Bellman recursion (§4.3), and the consequences and outer policy-iteration algorithm that consumes the gradient (§4.4).

4.2 The envelope identity

Let $\pi^\star(\boldsymbol{c}) = \arg\max_{\pi\in\Pi} V(\boldsymbol{c}, \pi)$ be the optimal policy at hardware $\boldsymbol{c}$. The joint-optimal geometry solves

$$\boldsymbol{c}^\star = \arg\max_{\boldsymbol{c}} V^\star(\boldsymbol{c}), \qquad V^\star(\boldsymbol{c}) \coloneqq V(\boldsymbol{c}, \pi^\star(\boldsymbol{c})). \tag{19}$$

A naive outer gradient requires differentiating through the policy argmax:

$$\frac{dV^\star}{d\boldsymbol{c}} = \left.\frac{\partial V}{\partial \boldsymbol{c}}\right|_{\pi=\pi^\star(\boldsymbol{c})} + \underbrace{\left.\frac{\partial V}{\partial \pi}\right|_{\pi=\pi^\star(\boldsymbol{c})}}_{=0 \text{ at optimum}} \cdot \frac{d\pi^\star}{d\boldsymbol{c}}. \tag{20}$$

The second term vanishes at the optimum by the first-order optimality condition (or, in discrete-action settings, by the envelope/Danskin theorem [49, 62, 92]). Consequently the outer gradient is simply

$$\boxed{\frac{dV^\star}{d\boldsymbol{c}} = \left.\frac{\partial V(\boldsymbol{c}, \pi^\star(\boldsymbol{c}))}{\partial \boldsymbol{c}}\right|_{\pi^\star \text{ held fixed}}.} \tag{21}$$

The policy Jacobian $d\pi^\star/d\boldsymbol{c}$ does not appear.

4.3 Unfolding the envelope gradient through the Bellman recursion

Equation (21) gives the envelope gradient in its abstract form. Making it operational requires unfolding through the Bellman recursion of Equations (4)–(5) to identify the concrete paths by which $\boldsymbol{c}$ enters $V^\star$. The result is a backward-recursion formula for $\partial V_k/\partial \boldsymbol{c}$ that mirrors the Bellman recursion for V_k, and that can be computed by reverse-mode automatic differentiation with a single structural step: freeze the argmax.

Define the Q-function

$$Q_k(\boldsymbol{b}, \boldsymbol{s}_k;\, \boldsymbol{c}) \;\coloneqq\; \mathbb{E}_{\boldsymbol{y}_k|\boldsymbol{b},\boldsymbol{s}_k,\boldsymbol{c}}\big[V_{k+1}\big(T(\boldsymbol{b}, \boldsymbol{s}_k, \boldsymbol{y}_k;\, \boldsymbol{c});\, \boldsymbol{c}\big)\big], \tag{22}$$

so that the Bellman recursion reads $V_k(\boldsymbol{b};\, \boldsymbol{c}) = \max_{\boldsymbol{s}_k} Q_k(\boldsymbol{b}, \boldsymbol{s}_k;\, \boldsymbol{c})$ with $\boldsymbol{s}_k^\star(\boldsymbol{b};\, \boldsymbol{c}) \coloneqq \arg\max_{\boldsymbol{s}_k} Q_k$. The total derivative of V_k with respect to $\boldsymbol{c}$ splits into an indirect piece (via $\boldsymbol{s}_k^\star$) and a direct piece:

$$\frac{\partial V_k}{\partial \boldsymbol{c}}(\boldsymbol{b};\, \boldsymbol{c}) \;=\; \underbrace{\frac{\partial Q_k}{\partial \boldsymbol{s}_k}(\boldsymbol{b}, \boldsymbol{s}_k^\star;\, \boldsymbol{c}) \cdot \frac{\partial \boldsymbol{s}_k^\star}{\partial \boldsymbol{c}}(\boldsymbol{b};\, \boldsymbol{c})}_{=0 \text{ at the argmax}} + \frac{\partial Q_k}{\partial \boldsymbol{c}}(\boldsymbol{b}, \boldsymbol{s}_k^\star;\, \boldsymbol{c}). \tag{23}$$

The first term vanishes at an interior argmax by first-order optimality ($\partial Q_k/\partial \boldsymbol{s}_k = 0$ at $\boldsymbol{s}_k^\star$). What survives is the partial derivative of Q_k holding $\boldsymbol{s}_k^\star$ fixed — the per-step envelope identity

$$\frac{\partial V_k}{\partial \boldsymbol{c}}(\boldsymbol{b};\, \boldsymbol{c}) \;=\; \frac{\partial Q_k}{\partial \boldsymbol{c}}(\boldsymbol{b}, \boldsymbol{s}_k^\star;\, \boldsymbol{c}). \tag{24}$$

Unfolding $\partial Q_k/\partial \boldsymbol{c}$ using the product and chain rules, three contributions enter:

$$\begin{aligned}\frac{\partial V_k}{\partial \boldsymbol{c}}(\boldsymbol{b};\, \boldsymbol{c}) \;=\; \mathbb{E}_{\boldsymbol{y}_k|\boldsymbol{b},\boldsymbol{s}_k^\star,\boldsymbol{c}}\Bigg[& \underbrace{\frac{\partial \log p(\boldsymbol{y}_k \mid \boldsymbol{b}, \boldsymbol{s}_k^\star, \boldsymbol{c})}{\partial \boldsymbol{c}}\, V_{k+1}\big(T(\cdots);\, \boldsymbol{c}\big)}_{\text{(i) score / outcome-likelihood term}} \\ & + \underbrace{\frac{\partial V_{k+1}}{\partial \boldsymbol{b}}\big(T;\, \boldsymbol{c}\big) \cdot \frac{\partial T(\boldsymbol{b}, \boldsymbol{s}_k^\star, \boldsymbol{y}_k;\, \boldsymbol{c})}{\partial \boldsymbol{c}}}_{\text{(ii) belief-update term}} \\ & + \underbrace{\frac{\partial V_{k+1}}{\partial \boldsymbol{c}}\big(T;\, \boldsymbol{c}\big)}_{\text{(iii) recursive term}} \Bigg],\end{aligned} \tag{25}$$

with terminal condition $\partial V_K/\partial \boldsymbol{c} = \partial \Phi_0(\boldsymbol{x}, \boldsymbol{b}_K;\, \boldsymbol{c})/\partial \boldsymbol{c}$ at the terminal belief. The three terms pick up the three ways $\boldsymbol{c}$ enters the Bellman backup once the argmax is frozen: through the outcome distribution (i), through the Bayesian belief update T (ii), and through the recursive dependence of V_{k+1} itself (iii).

Equation (25) mirrors Equation (5): same DP graph, same backward sweep order, same memoization structure. If V_k fits in memory — whether as an exact hash table, a count-tuple dictionary, or a Gaussian-moment parametrization — so does $\partial V_k/\partial \boldsymbol{c}$.

The operational recipe: argmax-freeze, then reverse-mode AD. One does not implement Equation (25) by hand. The workflow is:

1. **Forward pass.** Run backward induction from $k = K$ down to $k = 0$, populating V_k and recording $\boldsymbol{s}_k^\star$ at every reachable belief summary. No AD tape is active.
2. **Gradient pass.** Re-evaluate $V_0(p_0; \boldsymbol{c})$ by substituting the recorded $\boldsymbol{s}_k^\star$ in place of $\max_{\boldsymbol{s}_k}$. The resulting expression is argmax-free and, by Equation (21), identically equal to $V^\star(\boldsymbol{c})$.
3. **Reverse-mode AD.** Apply an AD system to the gradient-pass code. Because the expression is composed of closed-form primitives in $\boldsymbol{c}$ (the likelihood, the belief update, the terminal reward) and $\boldsymbol{c}$-free combinators (sums, products, argmax lookups), AD returns $\partial V_0/\partial \boldsymbol{c}$ exactly with no custom rule for the Bellman max.

Freezing the argmax turns the non-smooth max-containing recursion into an ordinary differentiable expression — that is the operational content of the envelope theorem. Without it, AD would try to differentiate through the max and produce subgradient terms that Equation (21) tells us are already zero at the optimum.

Concrete realizations across case studies. In Case Study A (radar, discrete belief), the per-step envelope gradient reduces to a finite sum over reachable beliefs; $\partial \Pr/\partial \boldsymbol{c}$ and $\partial T/\partial \boldsymbol{c}$ have closed forms in the beam-pattern parameter. In Case Study B (qubit, count-tuple state), the only $\boldsymbol{c}$-carrying primitive in the entire DP is the Ramsey likelihood of Equation (28); all of the count-tuple combinatorics, action grid, and memoization are $\boldsymbol{c}$-free, so the reverse-mode AD graph consists of one closed-form function composed through the DP tree via the chain rule. In Case Study C (photonic, Gaussian-posterior surrogate), the envelope theorem specializes to the implicit-function-theorem adjoint over the inner sensor-parameter argmax, and the per-step $\partial Q_k/\partial \boldsymbol{c}$ becomes a closed-form expression in the FDFD S-matrix derivatives; see §8.

4.4 Consequences and the outer policy-iteration algorithm

Equations (21) and (25) are the workhorses of the paper; together they give the sharp-max differentiable-DP outer gradient. Four consequences:

1. The outer gradient requires solving the inner problem once (to obtain $\pi^\star(\boldsymbol{c})$) and then differentiating V through the likelihood with the policy held fixed. This decouples the computational cost of the Bellman DP from the cost of the outer gradient.
2. When $\pi^\star(\boldsymbol{c})$ is computed only approximately (by deep RL, by an approximate value iteration, by a myopic one-step-lookahead), the envelope identity still says that at the *approximate* optimum the approximate policy Jacobian's contribution to the outer gradient is first-order small. The framework accommodates approximate inner solvers without bias.
3. The outer gradient is differentiable in $\boldsymbol{c}$ whenever the likelihood is. For photonic problems this is automatic via differentiable electromagnetic simulation [93, 94]; for quantum-circuit problems, via differentiable Hamiltonian dynamics [95].
4. A trajectory-sampling Monte-Carlo realization is available: sample trajectories under $\pi^\star$, differentiate the trajectory log-likelihood times the terminal reward under $\boldsymbol{c}$, average. This is the pathwise REINFORCE estimator [96, 97] specialized to the fixed-policy case.

The natural outer-loop algorithm alternates Bellman solves with these gradient steps:

Joint $(\boldsymbol{c}, \pi)$ policy iteration:

1. Initialize $\boldsymbol{c}^{(0)}$.
2. For $t = 1, 2, \ldots$ until convergence:
   - Inner: solve $\pi^\star(\boldsymbol{c}^{(t-1)})$ (exact DP, deep RL, myopic, etc.).
   - Outer: $\boldsymbol{c}^{(t)} = \boldsymbol{c}^{(t-1)} + \eta \cdot \frac{\partial V(\boldsymbol{c}^{(t-1)}, \pi^\star(\boldsymbol{c}^{(t-1)}))}{\partial \boldsymbol{c}}$ with step size η.
   - Optionally: box-project $\boldsymbol{c}^{(t)}$ onto the feasible set.

Convergence to a stationary point of $V^\star$ follows from standard policy-iteration theory [98, 99] when the inner solver is exact; the envelope theorem extends this to approximate inner solvers provided the approximation error is $o(\eta)$.

4.5 A note on continuous actions

Nothing in the differentiable-DP construction presupposes discrete actions. For continuous $\boldsymbol{s}$, the Bellman backup's inner maximization is a continuous optimization:

$$V_k(\boldsymbol{b}) = \max_{\boldsymbol{s}\in\mathcal{S}} \mathbb{E}_{\boldsymbol{y}}[V_{k+1}(T(\boldsymbol{b}, \boldsymbol{s}, \boldsymbol{y}))], \tag{26}$$

with the terminal condition V_K given by the inner expression of Φ at the terminal belief (Table 1). The $k = 0$ value under the optimal policy coincides with $V(\boldsymbol{c}, \pi^\star)$ of Equation (2). Standard tools, differential dynamic programming [100, 101], sample-based tree search [102, 103], local gradient inner optimization, HJB-PDE solutions [104], apply. The one caveat specific to our setting is that the standard count-statistic belief memoization used by discrete-action Bellman DP (Section 6) fails for continuous actions because distinct inner-optimization candidates produce distinct child beliefs. We handle this in Case Study C by using the Gaussian-posterior surrogate, in which the belief state is finite-dimensional and the Bellman backup admits a one-step-lookahead closed form; and in Case Study B by log-spaced discretization of the continuous Ramsey-delay dimension at $J = 10$ grid points.

5 A hierarchy of relaxations

The joint $(\boldsymbol{c}, \pi)$ optimization of Section 2 is exactly solvable only for small discrete POMDPs. When exact Bellman DP is intractable, two distinct philosophies emerge for handling the inner optimization.

Dichotomy at the exact-DP branch point. The dominant approach in modern reinforcement learning is to parametrize the policy π_θ (or the action-value Q_θ) with a neural network and optimize it by policy gradient, Q-learning, or actor–critic [105, 106]. *We do not pursue this branch in the present paper, though we do not rule it out for future work.* The reason is a well-documented pathology: an expressive NN policy can absorb the training signal and compensate for a poor $\boldsymbol{c}$ by investing in policy complexity, leaving $\boldsymbol{c}$ near its initialization. The outer loop converges to a *hardware-light* fixed point that looks good in simulation but relies on a near-trivial $\boldsymbol{c}$ hard to justify as a physical design. Known countermeasures — capacity control on π_θ (bottleneck width, weight decay, dropout), two-timescale updates with $\boldsymbol{c}$ stepped on a faster clock than θ, warm-starts from one of the prescribed inner rules below (rungs 3–4), and auxiliary regularizers penalizing hardware-null geometries — partially mitigate the risk but do not yet admit the clean sharp-max differentiable-DP gradient our framework exploits. We flag the combination of these remedies as a plausible path to extending joint-$(\boldsymbol{c}, \pi)$ to expressive NN policies in future work, and correspondingly exclude NN-parametrized policies from the hierarchy below on a *provisional* rather than principled basis.

The plausible-policy cascade (our path). The alternative, which we pursue, is to *prescribe* π in closed form as a *plausibly Bellman-optimal* rule at the affordable level of fidelity, and to focus the optimization machinery on $\boldsymbol{c}$. Each rule is a closed-form arg max over sensor parameters given a belief summary, chosen as a principled heuristic that approximates the Bellman optimum at its feasibility ceiling; the outer loop then moves $\boldsymbol{c}$. The hierarchy is a *cascade*: each rung is triggered by the previous becoming intractable.

Because the policy at every rung is prescribed as an explicit function of the current posterior (or one of its coarsenings), this philosophy *requires* an explicit Bayesian filter as the inference

primitive — the filter supplies the sufficient statistic the prescribed rule consumes. Explicit inference is therefore a *consequence* of prescribing a plausible policy, not an independent design choice. This is what the cascade shares with exact DP (rung 1) and what sharply distinguishes our branch from the NN branch above, where a learned encoder over the raw observation history replaces the Bayesian filter entirely. What descends the cascade is (i) the policy's proximity to the Bellman optimum — Bellman-optimal at rung 1, clairvoyant-CE at rung 2, myopic one-step at rungs 3–4 — and correspondingly (ii) the coarseness of the belief summary the policy consumes (and the filter maintains) — exact posterior, point estimate, full distribution, Gaussian moments. At no rung is the inference primitive replaced by a learned representation; the Bayesian filter is a first-class object throughout, because the prescribed policy requires it to be. The differentiable-DP gradient of Section 4 relies on this: it applies whenever the inner rule is a prescribed arg max over a belief-summary sufficient statistic, regardless of how coarsely that summary approximates the true posterior.

Rung 1 — Exact Bellman DP. $\pi_k^\star = \arg\max_{\boldsymbol{s}_k} \mathcal{Q}_k^\star(\boldsymbol{b}_k, \boldsymbol{s}_k)$, the true optimum. Inference: exact belief $\boldsymbol{b}_k$ (tabulated for discrete state; on a K_Φ-point grid for continuous state). Feasibility ceiling: reachable-belief memoization $\sim |\mathcal{S}|^K$ up to $\sim 10^6$ on a single node. When this is violated, move to rung 2.

Rung 2 — Certainty-equivalent control, full horizon. At step k, form a point estimate of the unknown state (posterior mean $\hat{\boldsymbol{x}}_k = \mathbb{E}[\boldsymbol{x} \mid \boldsymbol{b}_k]$ or posterior mode / MAP) and play the schedule the clairvoyant oracle would choose against $\hat{\boldsymbol{x}}_k$. Inference: explicit point estimate from the filter. Near-optimal when the posterior concentrates rapidly and the oracle map is Lipschitz in $\boldsymbol{x}$; known failure mode: under-exploration when $\hat{\boldsymbol{x}}_k$ is wrong. Feasibility ceiling: one oracle solve per decision step (itself an N-step enumeration or continuous optimization). When the per-step oracle solve is too expensive, move to rung 3.

Rung 3 — Myopic one-step information criterion. Collapse the look-ahead to one step: $\boldsymbol{s}_k = \arg\max_{\boldsymbol{s}_k} \mathbb{E}_{\boldsymbol{y}_k|\boldsymbol{b}_k,\boldsymbol{s}_k}\big[\Phi(\boldsymbol{x}, \boldsymbol{s}_k, \boldsymbol{c})\big]$, i.e., maximize the single-step expected information (or reward) gain conditional on the full current belief. Inference: full belief $\boldsymbol{b}_k$ maintained by exact Bayesian update. Near-optimal by submodularity bounds for information-theoretic rewards [107]; failure mode: myopia proper, i.e., no credit for information that pays off two or more steps out. Feasibility ceiling: one arg max over $\boldsymbol{s}_k$ plus a full-belief expectation per step. When the full-belief expectation is itself too expensive (high-dimensional continuous $\boldsymbol{x}$, for example), move to rung 4.

Rung 4 — Fisher-trace surrogate with EKF. Replace the information criterion by the Fisher-information trace evaluated at the posterior mean, and replace the exact Bayesian update by an extended Kalman filter: $\boldsymbol{s}_k = \arg\max_{\boldsymbol{s}_k} \operatorname{tr} J(\bar{\boldsymbol{x}}_k, \boldsymbol{s}_k, \boldsymbol{c})$ with $\boldsymbol{b}_k \approx \mathcal{N}(\bar{\boldsymbol{x}}_k, \Sigma_k)$ propagated by EKF. Inference: Gaussian moments $(\bar{\boldsymbol{x}}_k, \Sigma_k)$. This is the cheapest rung still compatible with continuous $\boldsymbol{c}$, continuous $\boldsymbol{x}$, and arbitrary state dimension. Failure modes: Gaussian approximation breaks near prior boundaries or under multimodal posteriors; EKF linearization fails when the forward model is strongly non-linear in $\boldsymbol{x}$ over the posterior support.

rung	inner policy rule	inference primitive	triggered when rung above is intractable
1	exact Bellman $\arg\max$	exact belief (grid / tabulated)	—
2	certainty-equivalent oracle at $\hat{\boldsymbol{x}}_k$	point estimate $\hat{\boldsymbol{x}}_k$	memoization $\lvert\mathcal{S}\rvert^K \gtrsim 10^6$
3	myopic $\arg\max_{\boldsymbol{s}_k} \mathbb{E}[\Phi \mid \boldsymbol{b}_k]$	full belief $\boldsymbol{b}_k$	per-step oracle too expensive
4	$\operatorname{tr} J(\bar{\boldsymbol{x}}_k, \boldsymbol{s}_k, \boldsymbol{c})$ + EKF	Gaussian $(\bar{\boldsymbol{x}}_k, \Sigma_k)$	full-belief expectation too expensive

Across all four rungs, the sharp-max differentiable-DP gradient applies cleanly to the outer $\boldsymbol{c}$-gradient at the inner $\arg\max$ (with the envelope theorem certifying exactness at the prescribed argmax), because the inner rule is a prescribed $\arg\max$ over a belief-summary maintained by an explicit Bayesian filter. The case studies that follow operate at rungs 1, 1, and 4 respectively: Case Study A (radar) at rung 1 with discrete state; Case Study B (scqubit) at rung 1 with a grid-discretized continuous state; Case Study C (photonic metasensor) at rung 4 with $\sim 10^5$ continuous hardware parameters. No single rung uniformly dominates. Each names its feasibility ceiling and failure mode explicitly so that a practitioner can descend the cascade deliberately, stopping at the first rung that is tractable for the problem in front of them.

An invitation to a broader research program. Once exact Bellman DP is relaxed, the inner policy becomes an approximation, and the full zoo of reinforcement-learning and prescribed-policy frameworks opens up — policy gradient, actor–critic, off-policy Q-learning, POMCP, information-directed sampling, distributional RL, Thompson sampling, model-based planning with learned dynamics, and many more — each with its own strengths, failure modes, and regimes of applicability. To our knowledge, this paper is the first to embed the optimization of continuous physical hardware $\boldsymbol{c}$ into this zoo, via sharp-max differentiable DP (the envelope-theorem-certified outer gradient) plus a prescribed, plausibly-Bellman-optimal inner rule. Correspondingly, each existing RL or prescribed-policy framework invites a $\boldsymbol{c}$-aware revision — the "what does this framework look like when the hardware is a co-equal decision variable?" question — whose answer will generally require new theory (ensuring the $\boldsymbol{c}$-gradient is well-defined and free of the hardware-light pathology), new algorithms (coordinating $\boldsymbol{c}$-updates with the inner policy on compatible timescales, integrating regularizers that penalize hardware-null geometries), and new numerics (quantifying how much each framework gains or loses when $\boldsymbol{c}$ is co-optimized rather than fixed). We view this as a new research area opened by the present formulation rather than exhausted by it; the four rungs above are the first systematic traversal, and many more remain.

A deeper stake: intelligence migrating from digital to physical. Information loss at the analog-to-digital boundary upper-bounds any downstream reconstruction; joint $(\boldsymbol{c}, \pi)$ optimization responds by shaping the physical geometry to participate in the inference rather than merely transduce for it. The relaxation cascade above carries this principle from the pedagogical small to the physically large. A co-optimized $\boldsymbol{c}$ is no longer a lens, an antenna, or a SQUID loop that merely transduces; it is a purpose-built information-preserving structure that filters, concentrates, and encodes the incoming signal in the geometry's own complex degrees of freedom, shaped by the task and by the policy that will consume its output. The broader research program we invite is, in this view, a piece of the larger migration of intelligence from digital to physical: in place of a generic measurement whose deficit is rescued by deep networks downstream, the analog structure carries its task-aware share, and the juncture of information loss is relocated, narrowed, and in some regimes transcended. What exact Bellman DP proves

possible in the pedagogical small, the relaxed rungs make feasible in the physically large — and the physical structure, once co-optimized, is no longer a conduit for intelligence but a participant in it.

6 Case Study A: Radar beam search (exact DP, discrete state, Φ_{MI})

The radar beam-search case is the smallest non-trivial POMDP that exhibits the full V_{oracle} / V_{fixed} / V_{adaptive} / $\mathbb{E}[\mathrm{IG}]$ decomposition, and it serves as the pedagogical entry point for the rest of the paper.

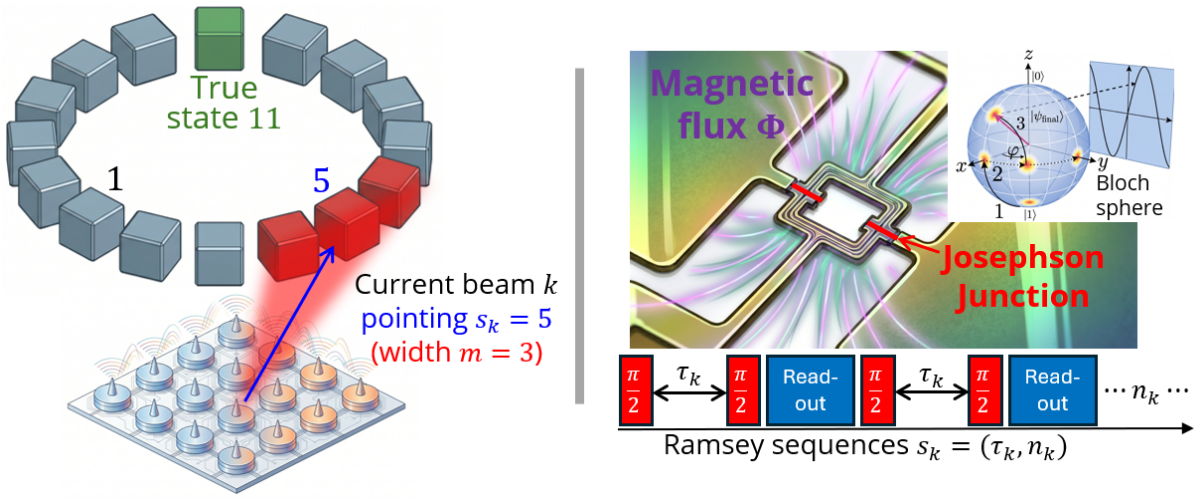

Figure 1: Scene illustrations for the two case studies. *Left* (Case Study A, §6): a phased-array antenna interrogates a ring of 16 target cells; the true state (green, cell 11), the current beam pointing at cell $\boldsymbol{s}_k = 5$ with width $m = 3$, and the adaptive redirect selected by $\pi^\star$ at the next epoch. *Right* (Case Study B, §7): the x-mon transmon cross on sapphire with a SQUID loop (two Josephson junctions) at the end of the $+x$ arm, the external flux Φ_{ext} threading the loop, a meandering readout resonator, and the Bloch-sphere view of the qubit; the bottom timeline shows the adaptive Ramsey sequences $\boldsymbol{s}_k = (\tau_k, n_k)$ chosen by $\pi^\star(\boldsymbol{b}_k)$ at each epoch.

6.1 Problem setup

State. $\boldsymbol{x} \in \{1, \ldots, |\mathcal{X}|\}$ is a target cell on a ring of $|\mathcal{X}| = 16$ cells. Prior $p_0(\boldsymbol{x}) = 1/|\mathcal{X}|$.

Action. At each of $K = 4$ epochs the sensor points a beam at cell $\boldsymbol{s}_k \in \{1, \ldots, |\mathcal{X}|\}$.

Observation. $\boldsymbol{y}_k \in \{0, 1\}$ (miss/detection) with Bernoulli likelihood

$$p(\boldsymbol{y}_k = 1 \mid \boldsymbol{x}, \boldsymbol{s}_k; \boldsymbol{c}) = p_{\mathrm{FA}} + (p_{\mathrm{D}} - p_{\mathrm{FA}}) \cdot G_{\boldsymbol{c}}(\theta_{\boldsymbol{s}_k} - \varphi_{\boldsymbol{x}}), \tag{27}$$

with $p_{\mathrm{D}} = 0.9$, $p_{\mathrm{FA}} = 0.05$, cell bearings $\varphi_{\boldsymbol{x}} = 2\pi(\boldsymbol{x} - 1)/|\mathcal{X}|$, and pointing angles $\theta_{\boldsymbol{s}_k}$. The geometry family is a top-hat (uniform-sensitivity, rectangular-profile) beam parametrized by integer width $m \in \{1, \ldots, |\mathcal{X}|\}$: the beam centered on cell a covers all cells within a ring distance of $\lfloor m/2 \rfloor$ from a, giving detection probability p_{D} uniformly inside the beam and p_{FA} uniformly outside. m is the discrete design variable we sweep below.

Reward. $\Phi = \Phi_{\mathrm{MI}}$, since the state is discrete and there is no natural MSE.

6.2 Enumeration results

The exact enumeration of V_{oracle}, V_{fixed}, and backward-induction V_{adaptive} at $|\mathcal{X}| = 16$, $K = 4$ is computationally modest (order 10^7 elementary operations, Julia runtime ~ 0.1 s on a single thread). In the narrow-beam ($m = 1$) and wide-beam ($m = 8$) regimes we obtain:

	narrow ($m = 1$)	wide ($m = 8$)
V_{oracle}	2.580 nats	1.923 nats
V_{adaptive}	0.619	1.700
V_{fixed}	0.499	1.411
$\mathbb{E}[\text{IG}]$	2.081	0.513
realized gain $V_{\text{adaptive}} - V_{\text{fixed}}$	0.121	0.289
EVPI saturation	5.8%	56.4%

Note the non-monotone relationship: the narrow regime has the largest $\mathbb{E}[\text{IG}]$ (more than $4\times$ the wide regime's) but the smallest realized adaptive gain (less than half). The narrow EVPI is "loose", most of the adaptive headroom is unreachable by any policy at $K = 4$, while the wide EVPI is relatively tight.

6.3 Sweep over the top-hat beam width

Sweeping the integer beam width $m \in \{1, \ldots, |\mathcal{X}|\}$ and reporting all four functionals (consecutive pairs of m produce identical beams under the integer-width parametrization, so the table lists one representative per degenerate pair; the endpoints $\{1, 2\}$ and $\{15, 16\}$ are both shown to make the pairing explicit):

m	V_{oracle}	V_{adaptive}	V_{fixed}	$\mathbb{E}[\text{IG}]$	realized	sat. %
1	2.580	0.619	0.499	**2.081**	0.121	5.8
2	2.580	0.619	0.499	2.081	0.121	5.8
3	2.293	1.401	1.103	1.190	0.299	25.1
5	2.081	1.631	1.343	0.737	0.287	38.9
7	1.924	1.700	**1.411**	0.513	0.289	56.4
9	1.879	**1.704**	1.401	0.478	**0.303**	**63.5**
11	2.074	1.632	1.320	0.754	0.312	41.4
13	2.312	1.389	1.041	1.271	0.347	27.3
15	2.514	0.611	0.442	2.071	0.169	8.1
16	2.514	0.611	0.442	2.071	0.169	8.1

The optimum for V_{adaptive} is at $m^\star = 9$ ($V^\star_{\text{adaptive}} = 1.704$ nats, saturation 63.5%). V_{fixed} peaks one cell earlier at $m = 7$. Realized adaptive gain $V_{\text{adaptive}} - V_{\text{fixed}}$ peaks at $m = 13$ (0.347 nats). And the ignorance gap $\mathbb{E}[\text{IG}]$ peaks at the two *extremes* $m = 1$ (narrow) and $m = 16$ (isotropic), far from any V_{adaptive} or V_{fixed} maximum.

6.4 Three-way comparison: EVPI picks the wrong geometry

Top-hat three-way comparison.

$\boldsymbol{c}$	$\mathbb{E}[\mathrm{IG}]$	V_{adaptive}	V_{fixed}	realized	sat.
top-hat $m = 1$ (narrow extreme)	**2.081**	0.619	0.499	0.121	5.8%
top-hat $m = 16$ (isotropic extreme)	2.071	0.611	0.442	0.169	8.1%
top-hat $m = 9$ (optimum)	0.478	**1.704**	**1.401**	**0.303**	**63.5%**

The optimum $m^\star = 9$ dominates both extremes on every actionable column: V_{adaptive}, V_{fixed}, realized gain, and saturation. The *only* column on which it is not best is $\mathbb{E}[\mathrm{IG}]$, where both extremes beat it by a factor of four. An EVPI-only geometry selection would reject $m^\star = 9$ in favor of either $m = 1$ or $m = 16$, both of which realize V_{adaptive} that is nearly three times worse. Quantitatively, $V_{\mathrm{adaptive}}(m = 9)/V_{\mathrm{adaptive}}(m = 1) = 1.704/0.619 \approx 2.75$: the EVPI-argmax geometry is strictly suboptimal in attainable adaptive value by a factor of $\sim 2.75\times$ at this problem size.

The same observation underwrites the conceptual claim of Section 3: $\arg\max_{\boldsymbol{c}} \mathbb{E}[\mathrm{IG}](\boldsymbol{c}) \neq \arg\max_{\boldsymbol{c}} V_{\mathrm{adaptive}}(\boldsymbol{c})$, the disagreement is operationally large, and only the full joint solve picks the right geometry.

6.5 Optimal adaptive policy structure

The DP-optimal policy tree reveals structure beyond what the scalar values expose. In the wide ($m = 8$) regime, the depth-2 head of the tree is:

history	action
$\emptyset$	cell 1
$y_1 = 0$ (miss)	cell 4
$y_1 = 1$ (detect)	cell 13

A detection at cell 1 (which covers cells $\{14, 15, 16, 1, 2, 3, 4\}$) narrows the support to the upper ring half; step 2 at cell 13 (covering $\{10, \ldots, 16\}$) bisects that half. This is binary search conditioned on the observation, the hallmark of adaptive value that fixed schedules cannot replicate.

7 Case Study B: Superconducting-qubit flux sensor (exact DP, continuous state, Φ_{Var})

Case Study A is discrete-state and has no natural MSE metric; here we take the same exact-DP machinery to a continuous-state problem whose deployment metric is Bayesian MSE, and compare joint-DP co-design against the joint Bayesian Cramér–Rao fixed-schedule baseline that dominates the classical Fisher-information literature for scalar estimation. Joint-DP reduces deployed MSE by 11.3× at $z = +132\sigma$. Figure 1 (right) shows the sensor geometry and the adaptive Ramsey sequences.

7.1 Problem setup

We consider the frequency-tunable transmon flux sensor of Danilin, Nugent & Weides [40]. The unknown is the external magnetic flux $\boldsymbol{x} = \Phi_{\mathrm{ext}}$ threading a SQUID loop, with prior $\boldsymbol{x} \sim \mathrm{Uniform}[0, \phi_{\max}\Phi_0]$ at $\phi_{\max} = 0.1$ (the rationale for this operating point is discussed in

§7.6). Ramsey interferometry at a fixed operating flux $\Phi^\star$ returns a Bernoulli outcome with likelihood

$$p(\boldsymbol{y}=1 \mid \boldsymbol{x},\tau;\boldsymbol{c}) = \tfrac{1}{2} + \tfrac{1}{2}\,\mathrm{th}(\boldsymbol{x},\boldsymbol{c})\, e^{-A(\boldsymbol{x},\boldsymbol{c})\tau - B(\boldsymbol{x},\boldsymbol{c})^2\tau^2}\, \cos[\Delta\omega(\boldsymbol{x},\boldsymbol{c})\tau], \tag{28}$$

with all noise rates (A, B) and spectral coefficients ($\Delta\omega$) closed-form in the 7-dimensional parametric geometry $\boldsymbol{c} = (f_q^{\max}, E_C/h, \kappa, \Delta_{qr}, T, A_\Phi, A_{I_c})$ (see [40] for the complete rate formulas). The action at epoch k is the pair $\boldsymbol{s}_k = (\tau_k, n_k)$, with τ_k the Ramsey delay and n_k the number of repetitions (yielding a Binomial(n_k, p) observation).

Reward. $\Phi = \Phi_{\mathrm{Var}}$, motivated by the Bayesian MSE alignment discussed in Section 2.3.

7.2 Exact Bellman DP via count-statistic memoization

The Ramsey likelihood is invariant under permutation of measurements at the same (τ, n) tuple. Consequently the belief $\boldsymbol{b}_k$ at step k is a sufficient statistic of the cumulative counts $(N_j, M_j)_{j=1}^{J}$ where N_j is the total number of measurements made at delay τ_j across all past epochs and M_j the count of $\boldsymbol{y} = 1$ outcomes. This makes the belief state a finite tuple of integers, and the Bellman DP admits exact memoization on a count-tuple hash table. At $K = 4$, $J = 10$ (log-spaced τ in $[10, 320]$ ns), $L = 2$ (repetitions $n \in \{1, 10\}$), $K_\Phi = 256$ (belief grid resolution), the memo contains on the order of 10^7 nodes and the DP completes in minutes on a single thread.

7.3 Joint-DP co-design

We run the policy iteration of Section 4 with outer Adam updates at learning rate 5×10^{-4} and periodic policy reoptimization, on a physically realistic restricted geometry box ($f_q^{\max} \in [3, 12]$ GHz, $E_C/h \in [0.15, 0.4]$ GHz, $\kappa \in [0.1, 5]$ MHz, $\Delta_{qr} \in [0.8, 5]$ GHz), with the noise amplitudes and temperature held fixed at the operating values reported in [40] ($A_\Phi = A_{I_c} = 10^{-6}$, $T = 40$ mK). The outer Adam rescales each $\boldsymbol{c}$ component by its box width before stepping, which is necessary because the raw natural-unit ranges of the 7 components span 14 orders of magnitude. The best $\boldsymbol{c}$ encountered over the trajectory is

$$\boldsymbol{c}_1^\star \;\approx\; (f_q^{\max} = 8.99 \text{ GHz},\; E_C/h = 0.254 \text{ GHz},\; \kappa = 0.47 \text{ MHz},\; \Delta_{qr} = 2.02 \text{ GHz}), \tag{29}$$

with $V_{\mathrm{adaptive}}(\boldsymbol{c}_1^\star) = -\mathbb{E}[\mathrm{Var}(\boldsymbol{b}_K)] = -7.49 \times 10^{-5}\,\Phi_0^2$.

7.4 Joint PCRB baseline

The Bayesian Cramér–Rao bound [55, 56] gives an alternative joint optimization:

$$(\boldsymbol{c}_2^\star, \boldsymbol{s}_2^\star) \;=\; \underset{\boldsymbol{c},\boldsymbol{s}_{1:K}}{\arg\max}\; \log J_P(\boldsymbol{s}_{1:K}, \boldsymbol{c}), \qquad J_P = J_0 + \mathbb{E}_{\boldsymbol{x}}\Big[\textstyle\sum_{k=1}^{K} n_k\, J_F(\boldsymbol{x}, \tau_k, \boldsymbol{c})\Big], \tag{30}$$

where $J_F(\boldsymbol{x}, \tau, \boldsymbol{c}) = (\partial p/\partial \boldsymbol{x})^2/[p(1-p)]$ is the per-shot scalar Fisher information. This is a *single-level* optimization over $\boldsymbol{c}$ and a fixed schedule $\boldsymbol{s}_{1:K}$ (no Bellman, no belief tree). Running the analogous outer Adam with inner schedule enumeration at per-iteration reoptimization gives:

$$\boldsymbol{c}_2^\star \;\approx\; (f_q^{\max} = 9.62 \text{ GHz},\; E_C/h = 0.272 \text{ GHz},\; \kappa = 0.10 \text{ MHz},\; \Delta_{qr} = 3.20 \text{ GHz}), \tag{31}$$

with $\boldsymbol{s}_2^\star = (\tau_{10}, n_2)^4 = (320 \text{ ns}, 10)^4$. The PCRB-optimal geometry $\boldsymbol{c}_2^\star$ prefers a larger qubit frequency, a larger qubit-resonator detuning, and the minimum-feasible κ, whereas the joint-DP optimum $\boldsymbol{c}_1^\star$ settles closer to the experimental nominal value of each. The two optima differ materially inside the restricted box, so the deployed-MSE gap reported below is a full joint co-design comparison rather than a shared-geometry test of policy structure alone.

7.5 Deployed MSE comparison

We deploy each design with $n_{\mathrm{mc}} = 20{,}000$ Monte-Carlo trajectories at $K_\Phi = 256$ posterior resolution, paired common-seed $\boldsymbol{x}_{\mathrm{true}}$ draws. Each deployment uses the posterior-mean estimator. Results:

quantity	value
$\overline{\mathrm{MSE}}_1$ (joint DP)	$(7.49 \pm 0.24) \times 10^{-5}\ \Phi_0^2$
$\overline{\mathrm{MSE}}_2$ (PCRB baseline)	$(8.43 \pm 0.05) \times 10^{-4}\ \Phi_0^2$
$J_P^{-1}(\boldsymbol{c}_2^\star, \boldsymbol{s}_2^\star)$ (CRB)	$1.03 \times 10^{-9}\ \Phi_0^2$
ratio $\overline{\mathrm{MSE}}_2/\overline{\mathrm{MSE}}_1$	**11.27**
z-score	$\mathbf{+132\sigma}$

Joint-DP $(\boldsymbol{c}_1^\star, \pi_1^\star)$ beats the PCRB baseline $(\boldsymbol{c}_2^\star, \boldsymbol{s}_2^\star)$ by $11.3\times$ on deployed Bayesian MSE at $K = 4$, at $z = +132\sigma$.

The CRB lower bound $J_P^{-1} \approx 10^{-9}$ is five-to-six orders of magnitude below both empirical MSEs, a reminder that the asymptotic CRB is loose for finite-horizon biased posterior-mean estimators on bounded-support priors [81, 108]. Neither design approaches the CRB lower bound; the comparison between the two is nevertheless clean and statistically decisive.

7.6 Why $\phi_{\max} = 0.1$: aliasing, horizon, and PCRB validity

The prior width $\phi_{\max}$ interacts with the tractable horizon $K = 4$ through Ramsey aliasing. The likelihood (28) is periodic in $\Delta\omega(\boldsymbol{x})\,\tau$, so a prior that spans many Ramsey fringes yields a multimodal posterior that can be disambiguated only by measurements at sufficiently varied τ — to leading order, at least $\log_2(N_{\mathrm{fringes}})$ distinct delays are required. At $\phi_{\max} \to \Phi_0/2$ (uninformative prior) the support spans many fringes and four epochs are too few for either a fixed or an adaptive schedule to resolve them; both joint-DP and joint-PCRB converge to comparable aliasing-floor MSE and the adaptive margin shrinks. At $\phi_{\max}$ well below this the prior is narrow enough that aliasing is essentially absent, both methods operate far below the floor, and the joint-DP margin grows rapidly. Separately, the PCRB is a local/asymptotic bound that elides posterior multimodality, so at large $\phi_{\max}$ it is mis-specified as a baseline: the comparison would be doubly uninformative (joint-DP K-limited, joint-PCRB outside its regime of validity). Choosing $\phi_{\max} = 0.1$ sits just inside the regime where joint-DP's adaptive finesse becomes observable at $K = 4$ while keeping the PCRB baseline a valid reference — a representative rather than extremal test of co-design quality.

8 Case Study C: Photonic metasensor (Gaussian-trace surrogate, 10^5-dim continuous $\boldsymbol{c}$, $\Phi_{\log J}$)

The two preceding case studies used exact Bellman DP. The photonic case study operates on a $\sim 10^5$-pixel permittivity landscape where exact DP is infeasible, and demonstrates the relaxation-hierarchy rung 4: Gaussian-posterior belief updates via an extended Kalman filter, the Bayesian Fisher-information-matrix reward $\Phi_{\log J}$, and implicit-function-theorem differentiation through the inner optimization that selects the adaptive probe phases.

8.1 Problem setup

The device is a 2×2 lattice of identical 2D cross-waveguide scatterers, each cell a four-port cross with a central 300×300-pixel permittivity design region; all four cells share the same inverse-designed geometry, giving $\boldsymbol{c} \in [0,1]^{300\times 300}$ (equivalently $\boldsymbol{c} \in \mathbb{R}^{90{,}000}$) as the outer optimization

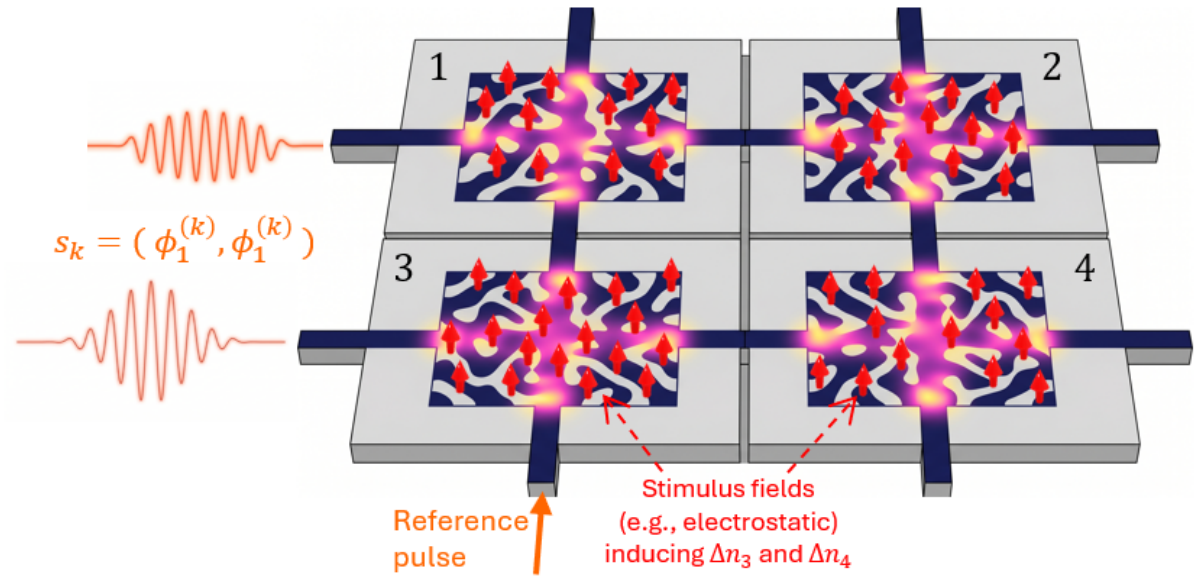

Figure 2: Case Study C overview. A 2×2 lattice of identical cross-waveguide scatterers (cells 1–4), each carrying the same inverse-designed binary-permittivity pattern in its central design region. Two coherent broadband light pulses enter at the left edge: an action pulse carrying the adaptive input phases $\boldsymbol{s}_k = (\phi_1^{(k)}, \phi_2^{(k)})$ at epoch k, and a fixed-phase reference pulse. External stimulus fields (e.g., electrostatic) acting on the cells induce per-cell refractive-index shifts $\Delta n_1, \ldots, \Delta n_4$, which constitute the hidden state $\boldsymbol{x} = \mathrm{vec}(\Delta n) \in \mathbb{R}^4$ to be estimated. The intensity overlays inside each cell visualize the electromagnetic near-field response under one such probe. Port powers measured at the 8 external output ports form the observation $\boldsymbol{y}_k \in \mathbb{R}^8$; an EKF reconstructs $\boldsymbol{x}$ across successive broadband measurement steps, with the adaptive policy re-selecting $\boldsymbol{s}_{k+1}$ from $\boldsymbol{x}$'s posterior after each step.

variable. The unknown state $\boldsymbol{x} = \mathrm{vec}(\Delta n) \in \mathbb{R}^4$ is the vector of per-cell refractive-index perturbations to be estimated. At each of $K_{\text{train}} = 3$ measurement epochs during co-design the sensor chooses a pair of input phases $\boldsymbol{s}_k = (\phi_1^{(k)}, \phi_2^{(k)})$ applied to two of the lattice's external input ports as coherent broadband drives across 20 frequencies; the observations are the $\boldsymbol{y}_k \in \mathbb{R}^8$ broadband port powers measured at the lattice's 8 external output ports, with additive Gaussian noise.

Reward. $\Phi = \Phi_{\log J}$ with J_N the Bayesian Fisher-information matrix accumulated over the K_{train} epochs. The Gaussian-posterior Kalman filter makes $\log\det J_N$ tractable in closed form at each epoch.

8.2 Physics and differentiability

The forward model is a 2D finite-difference frequency-domain (FDFD) Maxwell solver applied to a single unit cell, which returns one 4×4 S-matrix per frequency on a broadband sampling grid of 20 frequencies. We Taylor-expand the per-cell S-matrix to second order in the local refractive-index perturbation Δn,

$$S(\Delta n) \approx S_0 + \frac{\partial S}{\partial n}\Delta n + \frac{1}{2}\frac{\partial^2 S}{\partial n^2}\Delta n^2, \tag{32}$$

computed with a single FDFD solve per frequency and evaluated independently at each of the four cells with its own Δn_i. The four cells are wired together by the port-connection equations to assemble the lattice's 8-port S-matrix; the 8 external-port powers, summed over the 20

frequencies against a Gaussian spectral weight, form the observation $\boldsymbol{y} \in \mathbb{R}^8$. The Jacobian $\partial \boldsymbol{y}/\partial \boldsymbol{c}$ passes through the cell S-matrix coefficients via the FDFD adjoint method [93, 94].

8.3 Joint optimization

The outer loop is the policy iteration of Section 4. At each $\boldsymbol{c}$ the inner optimization $\boldsymbol{s}^\star(\boldsymbol{b}_k; \boldsymbol{c}) = \arg\max_{\boldsymbol{s}} \Phi(\boldsymbol{x}, \boldsymbol{s}, \boldsymbol{c})$ is solved by grid search followed by L-BFGS refinement. Its derivative with respect to $\boldsymbol{c}$, which the chain rule requires, is computed in closed form from the implicit function theorem applied at the inner argmax, avoiding nested automatic differentiation. To produce a geometry compatible with lithographic fabrication we apply a density filter to $\boldsymbol{c}$ (convolving with a radius-R kernel to enforce a minimum feature size) and then a tanh threshold projection with sharpness parameter β [109–112]. During optimization β is increased along a continuation schedule $16 \to 32 \to 64 \to 128 \to 256$ so that the projection gradually sharpens from a smooth blur to a near-binary mask.

8.4 Optimized binary design and loss trajectory

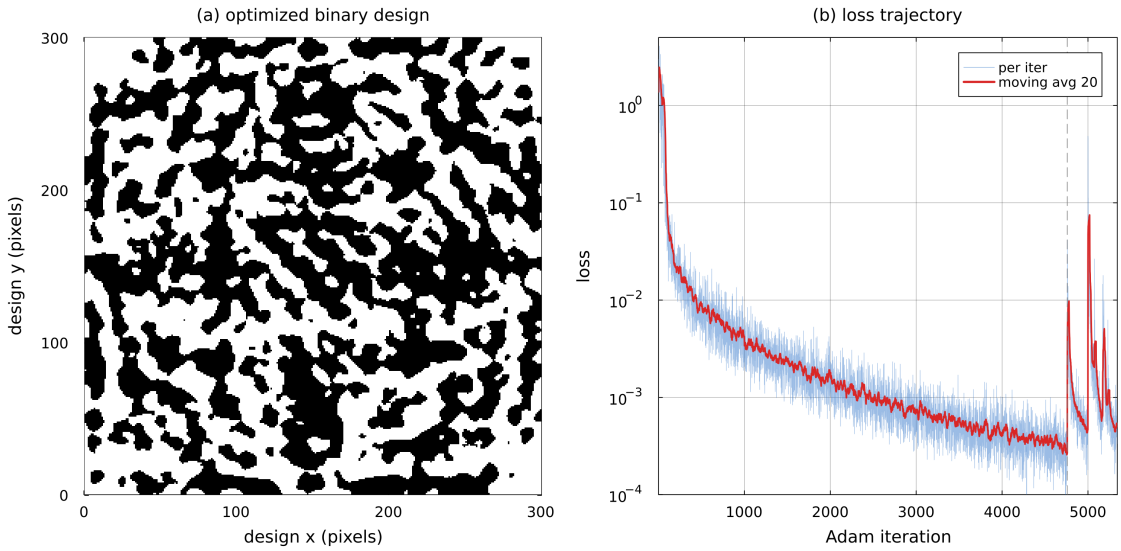

Figure 3: Final optimized binary permittivity of the 300×300 design region (main panel). Inset: joint-DP loss trajectory from a fully random initialization through the full co-design; light blue, per-iteration loss; red, 20-iteration moving average; dashed, the transition into the filter-constrained projection phase.

8.5 Monte-Carlo EKF deployment

Deploying the optimized binary design with the EKF reconstruction and fresh perturbation draws (200 episodes, $K_{\text{deploy}} = 10$ EKF steps per episode):

metric	optimized	random initial	ratio
EKF relative error (10 steps)	0.97%	120%	**123**× better
BFIM trace (step 1)	2.50×10^{13}	4.01×10^{10}	**622**×
Episode loss (10 steps)	3.5×10^{-4}	6.5	$\mathbf{1.84 \times 10^4}$×

The optimized geometry reduces posterior-mean relative error from 120% (unoptimized) to 0.97% (optimized) after ten EKF measurement steps, a 123× improvement, while delivering

622× more Fisher information per measurement. After a single broadband measurement the EKF reduces relative error from 184% (prior) to 3.4% (posterior), essentially solving the estimation problem with one observation. These ratios are enormous because the problem starts ill-conditioned: a random permittivity distribution produces a nearly non-invertible S-matrix. The improvement illustrates what high-dimensional $\boldsymbol{c}$-optimization buys when jointly designed with the EKF reconstruction.

8.6 Feasibility discussion

At the design frequency $\lambda = 1550$ nm, the shot-noise-limited detection threshold we report corresponds to $\sim 64\,\mu$W per port at $1\,\mu$s integration, within the operating range of standard telecom lasers. The final binary permittivity distribution is smooth at the density-filter length scale ($R = 5$ pixels ≈ 200 nm feature size at 5-nm/pixel lattice) and compatible with electron-beam lithography on standard SOI [113, 114].

9 Discussion

9.1 The common shape of the adaptive advantage

Across the three case studies the adaptive policy's margin over the best fixed-schedule baseline comes from one move: redirecting the next measurement based on whichever partition of the state space the current data has already ruled out. In Case A, a detection at cell 1 halves the ring's support before step 2 targets cell 13. In Case B, the first epoch at $\tau \approx 22$ ns collapses the multimodal posterior before later epochs switch to $\tau = 320$ ns to sharpen within a single fringe; the PCRB fixed schedule commits to $\tau = 320$ ns from step 1 and cannot undo the aliasing. In Case C, each new pair of probe phases is re-aligned with the largest remaining eigenvector of the EKF posterior covariance, so that successive measurements shrink the least-constrained direction of the state. Prior-averaged information criteria, whether mutual information, Fisher information, or EVPI, marginalize over the data before committing an action, so the geometry they recommend is tuned to whichever schedule is best on average rather than best conditionally. The joint-DP margin is therefore not a numerical artifact of any particular problem: it is the signature of a hardware operating point chosen in anticipation of a conditional policy, and hardware chosen without that anticipation cannot, in general, recover it by attaching an adaptive policy after the fact.

9.2 Why the PCRB baseline loses even when it co-optimizes $\boldsymbol{c}$

The 11.3× Case B margin sharpens this general statement. Both $\boldsymbol{c}_1^\star$ and $\boldsymbol{c}_2^\star$ are genuine numerical optima of their respective objectives inside the same physical box; nothing prevents the PCRB optimizer from landing on $\boldsymbol{c}_1^\star$ if that geometry gave it more Fisher information. It lands on a different point $\boldsymbol{c}_2^\star$ and still produces MSE an order of magnitude worse, because the two objectives reward different physical behavior. The Fisher information at long Ramsey delay is high because one long measurement gives sharp local phase resolution, but at prior width $\phi_{\max} = 0.1\,\Phi_0$ the prior support spans many Ramsey fringes and the single-level Fisher criterion has no machinery for disambiguating which fringe the true Φ_{ext} lies on. The Bellman-optimal policy does have that machinery: a short-τ first epoch reduces the multimodal posterior to a unimodal one, and the remaining epochs exploit long-τ resolution within the selected fringe. The PCRB-optimal $\boldsymbol{c}$ is therefore tuned to a schedule the adaptive policy would never deploy; and because the coherence and drive parameters that favor long-τ fringe visibility differ from those that favor short-τ aliasing-breaking, the two optimizers converge to different geometries entirely. The belief-conditional action selection is not a refinement one can bolt on afterwards, which

is what makes co-design, rather than serial "pick $\boldsymbol{c}$ then pick π" optimization, the minimum principled procedure.

9.3 What the EVPI *is* still good for

We advocate demoting the EVPI from a design target to a workflow tool. Three concrete roles survive:

1. A pre-screen when the full joint DP is computationally prohibitive but schedule enumeration remains tractable. $\mathbb{E}[\mathrm{IG}](\boldsymbol{c})$ requires only V_{oracle} and V_{fixed}, both computable by direct enumeration of candidate schedules, bypassing the backward-in-time value iteration that V_{adaptive} demands; the cheapness of EVPI is inherited from the cheapness of those enumerations (small discrete action sets or low-dimensional continuous actions on a reasonable grid), and evaporates when the schedule space is large enough that enumeration itself becomes the bottleneck.
2. A post-hoc diagnostic showing how much adaptive headroom is available at the joint-optimum geometry. Small $\mathbb{E}[\mathrm{IG}](\boldsymbol{c}^\star)$ means the joint-optimum is close to the oracle ceiling; large values indicate remaining value from longer horizons, richer action sets, or tighter priors.
3. An interpretive lens on the shape of the $\boldsymbol{c}$-landscape. The non-coincidence of $\arg\max \mathbb{E}[\mathrm{IG}]$ and $\arg\max V_{\text{adaptive}}$ encodes physics: extreme geometries have large ceilings that no policy can reach; interior geometries have smaller ceilings that adaptive policies can consume a larger fraction of.

9.4 Limitations

- **Horizon.** Case Studies A and B are at small horizon ($K = 4$) where exact DP is tractable. At larger horizons the Bellman memoization table blows up, forcing descent down rungs 2–4 of the hierarchy. The Gaussian-posterior EKF surrogate (rung 4) extrapolates to arbitrary K; its accuracy relative to exact DP at moderate K is an open question.
- **Parametric $\boldsymbol{c}$ in Case B.** Case Study B's 7-dimensional geometry is the closed-form-parametrized subset of the transmon circuit, not the full layout design space. Extending $\boldsymbol{c}$ to include SQUID-loop mutual inductances via analytically differentiable Biot–Savart [115] is architecturally supported and expected to give joint-DP further advantage; this is ongoing work.
- **Finite sample MSE bound.** The Bayesian CRB $1/J_P$ is an asymptotic lower bound. For biased Bayesian estimators on bounded priors at small K the bound is loose by orders of magnitude (Case Study B: $\sim 10^5$–$10^6\times$ gap between the CRB and the deployed MSE of either design). Tighter finite-sample bounds such as the Weiss–Weinstein [116], Bobrovsky–Zakai, and Ziv–Zakai [82] bounds would be more informative anchors; integrating them into the framework is a natural next step.
- **Continuous-action Bellman.** The count-statistic memoization exploited in Case Study B breaks for continuous actions. Practical paths include fine discretization (as we did), Gaussian belief approximations (rung 4), or particle-filter DP. A full continuous-action Bellman for the scqubit problem at $K \geq 4$ would require one of these paths; we presented the discretization path as the cheapest.

9.5 Open directions: dynamical state

All three case studies treat the unknown $\boldsymbol{x}$ as static over the K-step measurement horizon. The natural next frontier is *dynamical* $\boldsymbol{x}$: targets that move, circuit parameters that drift, fields that fluctuate across the sensor integration window. Three things change when the state evolves. First, the belief filter becomes a recursive predict-correct step rather than a single-shot

posterior update, and for non-Gaussian posteriors the inference primitive has to move from grids and EKFs to particle filters. Second, the Bellman backup has to anticipate state transitions over the K steps, not just measurement accumulation, so the dynamics time scale interacts with the horizon and with each rung's coarsening of the belief summary. Third, when the action space is high-dimensional or continuous (beam angles, integration times, waveform parameters) the prescribed closed-form arg max of §5 becomes unwieldy and a neural-network policy becomes the natural inner solver, which re-introduces the hardware-light pathology flagged at the branch point of the cascade. The sharp-max differentiable-DP outer gradient survives this transition unchanged; what does not survive is the exactness of the inner solve, and the clean cascade of rungs 1–4 is replaced by a spectrum of training regimes whose geometry gradients depend on policy capacity, dynamics time scale, and filter fidelity together.

We pursue this extension in a companion line of work on real-time tracking of a nonlinear dynamical system with a jointly-designed physical sensor and a learned measurement policy; full treatment is deferred to a future paper. The present work establishes the static-state foundation on which that extension stands.

10 Conclusion

For any sensor whose physical hardware is designed once but whose measurement policy runs for the lifetime of the device, programmable metasurfaces, cognitive radars, frequency-tunable quantum sensors, reconfigurable photonic-integrated meshes, joint optimization of geometry and policy is the minimum principled procedure. Sharp-max differentiable dynamic programming — with the envelope theorem as the exactness certificate — makes the joint optimization tractable; a principled relaxation hierarchy makes it portable across scales; a family of compatible scalar rewards makes it applicable across communities. The Expected Value of Perfect Information, classical and instructive, serves as a diagnostic: its argmax disagrees with the joint-DP argmax by factors of $\sim 3\times$ in attainable adaptive value. The deployment-aligned comparison of joint-DP and Bayesian-CRB co-design on a superconducting-qubit flux sensor gives the operational headline: an $11.3\times$ reduction in Bayesian MSE at $z = +132\sigma$, with both hardware geometries numerically co-optimized under their respective objectives.

Code and data. The Julia implementation of all three case studies, including the exact Bellman DP with count-statistic memoization (Case Studies A and B), the sharp-max differentiable-DP outer Adam with per-component box-width preconditioning, the MSE-terminal reward (`Bellman.jl`, `Gradient.jl`, `JointOpt.jl`), the joint PCRB baseline (`PCRB.jl`), and the photonic BFIM+EKF surrogate (`SimGeomBroadBand.jl`, `BFIMGaussian.jl`, `PhEnd2End.jl`), is publicly available. Numerical outputs of all tables in Sections 6–8 are serialized and shipped with the source; reproducing any table requires a single Julia invocation.

Acknowledgments. [redacted for blind review; funding, collaborators, and facility acknowledgments here.]

References

[1] Zin Lin and Steven G. Johnson. Overlapping domains for topology optimization of large-area metasurfaces. *Optics Express*, 27:32445–32461, 2019.

[2] Zin Lin, Victor Liu, Raphaël Pestourie, and Steven G. Johnson. Topology optimization of freeform large-area metasurfaces. *Optics Express*, 27:15765–15775, 2019.

[3] Sean Molesky, Zin Lin, Alexander Y. Piggott, Weiliang Jin, Jelena Vuckovic, and Alejandro W. Rodriguez. Inverse design in nanophotonics. *Nature Photonics*, 12:659–670, 2018.

[4] Jakob S. Jensen and Ole Sigmund. *Topology optimization for nano-photonics*. 2011. Laser & Photonics Reviews 5(2).

[5] Jonathan A. Fan et al. Freeform metasurface design with deep generative models. *Applied Physics Reviews*, 2020.

[6] Rasmus E. Christiansen, Zin Lin, et al. [check] fullwave freeform topology optimization review. 2021.

[7] Owen D. Miller, Yiwen Yang, Steven G. Johnson, et al. Fundamental limits to optical response in absorptive systems. *Optics Express*, 24:3329–3364, 2016.

[8] Owen D. Miller et al. Fundamental bounds on light-matter interaction. *[CHECK review]*, 2023.

[9] Vincent Sitzmann, Steven Diamond, Yifan Peng, Xiong Dun, Stephen Boyd, Wolfgang Heidrich, Felix Heide, and Gordon Wetzstein. End-to-end optimization of optics and image processing for achromatic extended depth of field and super-resolution imaging. *ACM Transactions on Graphics (SIGGRAPH)*, 2018.

[10] Julie Chang and Gordon Wetzstein. Deep optics for monocular depth estimation and 3d object detection. *ICCV*, 2019.

[11] Christopher A. Metzler, Hayato Ikoma, Yifan Peng, and Gordon Wetzstein. Deep optics for single-shot high-dynamic-range imaging. *CVPR*, 2020.

[12] Ethan Tseng, Shane Colburn, James Whitehead, Luocheng Huang, Seung-Hwan Baek, Arka Majumdar, and Felix Heide. Neural nano-optics for high-quality thin lens imaging. *Nature Communications*, 2021.

[13] Michael R. Kellman, Emrah Bostan, Nicole A. Repina, and Laura Waller. Physics-based learned design: optimized coded-illumination for quantitative phase imaging. *IEEE Transactions on Computational Imaging*, 5:344–353, 2019.

[14] Felix Heide, Markus Steinberger, Yun-Ta Tsai, Mushfiqur Rouf, Dawid Pajak, Dikpal Reddy, Orazio Gallo, Jing Liu, Wolfgang Heidrich, Karen Egiazarian, Jan Kautz, and Kari Pulli. FlexISP: A flexible camera image processing framework. *ACM Transactions on Graphics (SIGGRAPH Asia)*, 2014.

[15] Nick Antipa, Grace Kuo, Reinhard Heckel, Ben Mildenhall, Emrah Bostan, Ren Ng, and Laura Waller. DiffuserCam: lensless single-exposure 3d imaging. *Optica*, 2018.

[16] S. Bhargava et al. End-to-end topology-optimized imaging. *preprint*, 2024.

[17] Zin Lin, Charles Roques-Carmes, Raphaël Pestourie, Marin Soljačić, Arka Majumdar, and Steven G. Johnson. End-to-end nanophotonic inverse design for imaging and polarimetry. *Nanophotonics*, 10:1177–1187, 2021.

[18] Qian Wang et al. [check] reconfigurable metasurface inverse design. 2018.

[19] Leslie Pack Kaelbling, Michael L. Littman, and Anthony R. Cassandra. Planning and acting in partially observable stochastic domains. *Artificial Intelligence*, 101:99–134, 1998.

[20] Richard D. Smallwood and Edward J. Sondik. The optimal control of partially observable Markov processes over a finite horizon. *Operations Research*, 21(5):1071–1088, 1973.

[21] Edward J. Sondik. *The optimal control of partially observable Markov processes.* PhD thesis, Stanford University, 1971.

[22] Chris Kreucher, Keith Kastella, and Alfred O. Hero, III. Sensor management using an active sensing approach. *Signal Processing*, 85(3):607–624, 2005.

[23] Alfred O. Hero and Douglas Cochran. Sensor management: Past, present, and future. *IEEE Sensors Journal*, 11(12):3064–3075, 2011.

[24] A. O. Hero, D. A. Castañón, D. Cochran, and K. Kastella, editors. *Foundations and Applications of Sensor Management.* Springer, 2008.

[25] Yuxin Chen and Andreas Krause. Near-optimal bayesian active learning with correlated and noisy tests. In *AISTATS*, 2017.

[26] Andreas Krause and Carlos Guestrin. Near-optimal nonmyopic value of information in graphical models. In *UAI*, 2005.

[27] D. V. Lindley. On a measure of the information provided by an experiment. *Annals of Mathematical Statistics*, 27(4):986–1005, 1956.

[28] Xun Huan and Youssef M. Marzouk. Sequential Bayesian optimal experimental design via approximate dynamic programming. *arXiv:1604.08320*, 2016.

[29] Wanggang Shen and Xun Huan. Bayesian sequential optimal experimental design for nonlinear models using policy gradient reinforcement learning. *arXiv:2110.15335*, 2023.

[30] Alen Alexanderian. Optimal experimental design for infinite-dimensional Bayesian inverse problems governed by PDEs: a review. *Acta Numerica*, 2021.

[31] Kathryn Chaloner and Isabella Verdinelli. Bayesian experimental design: A review. *Statistical Science*, 10(3):273–304, 1995.

[32] Kristine L. Bell, Chris J. Baker, Graeme E. Smith, Joel T. Johnson, and Muralidhar Rangaswamy. Cognitive radar framework for target detection and tracking. *IEEE Journal of Selected Topics in Signal Processing*, 2015.

[33] Simon Haykin. Cognitive radar: a way of the future. *IEEE Signal Processing Magazine*, 23:30–40, 2006.

[34] Maria S. Greco, Fulvio Gini, Pietro Stinco, and Kristine L. Bell. Cognitive radars: On the road to reality. *IEEE Signal Processing Magazine*, 2018.

[35] Mourad Bouhou et al. POMDP-driven cognitive massive MIMO radar. *IEEE Transactions on Radar Systems*, 2025. arXiv:2410.17967.

[36] Vittorio Giovannetti, Seth Lloyd, and Lorenzo Maccone. Quantum metrology. *Physical Review Letters*, 96:010401, 2006.

[37] Samuel L. Braunstein and Carlton M. Caves. Statistical distance and the geometry of quantum states. *Physical Review Letters*, 72:3439–3443, 1994.

[38] Matteo G. A. Paris. Quantum estimation for quantum technology. *International Journal of Quantum Information*, 7:125–137, 2009.

[39] Carlton M. Caves. Quantum-mechanical noise in an interferometer. *Physical Review D*, 23:1693–1708, 1981.

[40] Sergey Danilin, Nicholas Nugent, and Martin Weides. Tunable transmon qubit as a flux sensor: joint device-protocol optimization. *arXiv:2211.08344v4*, 2024.

[41] A. Yu. Kitaev. Quantum measurements and the Abelian stabilizer problem. *arXiv:quant-ph/9511026*, 1995.

[42] B. Braverman, A. Kawasaki, E. Pedrozo-Peñafiel, S. Colombo, C. Shu, Z. Li, E. Mendez, M. Yamoah, L. Salvi, D. Akamatsu, et al. Near-unitary spin squeezing in ^{171}yb. *Physical Review Letters*, 122:223203, 2019.

[43] Okan Yurduseven, Jonah N. Gollub, Daniel L. Marks, and David R. Smith. Printed aperiodic cavity for computational and microwave imaging. *IEEE Microwave and Wireless Components Letters*, 2017.

[44] Mohammadreza F. Imani, Jonah N. Gollub, Okan Yurduseven, Aaron V. Diebold, Michael Boyarsky, Thomas Fromenteze, Laura Pulido-Mancera, Timothy Sleasman, and David R. Smith. Review of metasurface antennas for computational microwave imaging. *IEEE Transactions on Antennas and Propagation*, 2020.

[45] S. Pirandola et al. Advances in photonic quantum sensing. *Nature Photonics*, 12:724–733, 2018.

[46] C. L. Degen, F. Reinhard, and P. Cappellaro. Quantum sensing. *Reviews of Modern Physics*, 89:035002, 2017.

[47] Harald Cramér. *Mathematical Methods of Statistics*. Princeton University Press, 1946.

[48] C. R. Rao. Information and the accuracy attainable in the estimation of statistical parameters. *Bulletin of the Calcutta Mathematical Society*, 37:81–91, 1945.

[49] Paul Milgrom and Ilya Segal. Envelope theorems for arbitrary choice sets. *Econometrica*, 70(2):583–601, 2002.

[50] Lawrence M. Benveniste and Jose A. Scheinkman. On the differentiability of the value function in dynamic models of economics. *Econometrica*, 47(3):727–732, 1979.

[51] Arthur Mensch and Mathieu Blondel. Differentiable dynamic programming for structured prediction and attention. In *Proceedings of the 35th International Conference on Machine Learning (ICML)*, 2018.

[52] Brandon Amos, Ivan D. Jimenez Rodriguez, Jacob Sacks, Byron Boots, and J. Zico Kolter. Differentiable MPC for end-to-end planning and control. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.

[53] Justin Domke. Generic methods for optimization-based modeling. In *Proceedings of the 15th International Conference on Artificial Intelligence and Statistics (AISTATS)*, 2012.

[54] Mathieu Blondel, Quentin Berthet, Marco Cuturi, Roy Frostig, Stephan Hoyer, Felipe Llinares-López, Fabian Pedregosa, and Jean-Philippe Vert. Efficient and modular implicit differentiation. *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.

[55] Harry L. Van Trees and Kristine L. Bell. *Bayesian Bounds for Parameter Estimation and Nonlinear Filtering/Tracking*. Wiley-IEEE Press, 2007.

[56] Petr Tichavský, Carlos H. Muravchik, and Arye Nehorai. Posterior Cramér–Rao bounds for discrete-time nonlinear filtering. *IEEE Transactions on Signal Processing*, 46(5):1386–1396, 1998.

[57] Richard Bellman. *Dynamic Programming.* Princeton University Press, 1957.

[58] Michael R. Kellman, Emrah Bostan, Michael Chen, and Laura Waller. Data-driven design for Fourier ptychographic microscopy. *Optica*, 7:1075–1082, 2020.

[59] Alex Muthumbi, Amey Chaware, Kangwook Kim, Kevin C. Kim, Jaebum Park, Michael Chen, Laura Waller, and Roarke Horstmeyer. Learned sensing: jointly optimized microscope hardware for accurate image classification. *Biomedical Optics Express*, 10:6351–6369, 2019.

[60] Emrah Bostan, Reinhard Heckel, Michael Chen, Michael Kellman, and Laura Waller. Deep phase decoder: self-calibrating phase microscopy with an untrained deep neural network. *Optica*, 7:559–562, 2020.

[61] Michael R. Kellman, Nicole A. Repina, Emrah Bostan, and Laura Waller. [check] data-driven illumination design for defocus-robust imaging. 2020.

[62] John M. Danskin. The theory of max-min, with applications. *SIAM Journal on Applied Mathematics*, 14:641–664, 1966.

[63] Ronald A. Howard. Information value theory. *IEEE Transactions on Systems Science and Cybernetics*, 2(1):22–26, 1966.

[64] Howard Raiffa and Robert Schlaifer. *Applied Statistical Decision Theory.* Harvard Business School, 1961.

[65] Pengning Chao, Benjamin Strekha, others, Sean Molesky, and Alejandro W. Rodriguez. Maximum electromagnetic local density of states via material structuring: T-matrix bounds. *Nature Reviews Physics*, 2022.

[66] Sean Molesky, Prashanth S. Venkataram, Weiliang Jin, and Alejandro W. Rodriguez. Fundamental limits to radiative heat transfer: theory. *Physical Review B*, 2020.

[67] Sean Molesky, Pengning Chao, and Alejandro W. Rodriguez. Hierarchical mean-field t-matrix bounds on photonic scattering. *Physical Review Research*, 2020.

[68] Sean Molesky, Pengning Chao, Weiliang Jin, and Alejandro W. Rodriguez. T-operator limits on optical communication: metaoptics, computation, and input-output transformations. *Physical Review Research*, 2022.

[69] Zeyu Kuang and Owen D. Miller. Computational bounds to light-matter interactions via local conservation laws. *Physical Review Letters*, 2020.

[70] Benjamin Strekha, Pengning Chao, Sean Molesky, and Alejandro W. Rodriguez. [check] bounds on angle-dependent scattering via t-operator. 2022.

[71] Mats Gustafsson, Kurt Schab, Lukas Jelinek, and Miloslav Capek. Upper bounds on absorption and scattering. *New Journal of Physics*, 2020.

[72] Kurt Schab et al. [check] radiation-q bounds. 2020.

[73] Benjamin Strekha et al. [check] earlier bounds paper. 2020.

[74] Sean Molesky, Zin Lin, Alexander Y. Piggott, Weiliang Jin, Jelena Vuckovic, and Alejandro W. Rodriguez. Inverse design in nanophotonics. *Nature Photonics*, 12:659–670, 2018.

[75] Prashanth S. Venkataram, Sean Molesky, Weiliang Jin, and Alejandro W. Rodriguez. Fundamental limits to radiative heat transfer: the limited role of nanostructuring in the near-field. *Physical Review Letters*, 2020.

[76] Lukas Jelinek, Miloslav Capek, and Kurt Schab. [check]. 2021.

[77] Stefan Rotter and Sylvain Gigan. Light fields in complex media: Mesoscopic scattering meets wave control. *Reviews of Modern Physics*, 89:015005, 2017.

[78] Dorian Bouchet, Stefan Rotter, and Allard P. Mosk. Maximum information states for coherent scattering measurements. *Nature Physics*, 17:564–568, 2021.

[79] Hannah M. Carrington, Dorian Bouchet, and Stefan Rotter. Fisher-information and optimal probes for wave scattering. *[CHECK]*, 2020.

[80] Hugo Defienne, B. Ndagano, et al. [check] quantum imaging and fisher information. 2021.

[81] Mankei Tsang, Ranjith Nair, and Xiao-Ming Lu. Quantum theory of superresolution for two incoherent optical point sources. *Physical Review X*, 6:031033, 2016.

[82] Mankei Tsang. Ziv–zakai error bounds for quantum parameter estimation. *Physical Review Letters*, 108, 2012.

[83] Mankei Tsang. Resolving starlight: a quantum perspective. *Contemporary Physics*, 2019.

[84] Claude E. Shannon. A mathematical theory of communication. *Bell System Technical Journal*, 27:379–423, 623–656, 1948.

[85] Paola Sebastiani and Henry P. Wynn. Maximum entropy sampling and optimal Bayesian experimental design. *J. Royal Statistical Society Series B*, 62:145–157, 2000.

[86] A. W. van der Vaart. *Asymptotic Statistics*. Cambridge University Press, 2000.

[87] Allan [CHECK] Clark. [check] bernstein–von mises review. 2021.

[88] Luc Pronzato. Optimal experimental design for non-linear models. 2010.

[89] James O. Berger. *Statistical Decision Theory and Bayesian Analysis*. Springer, 1985.

[90] Mert Pilanci, Martin J. Wainwright, and Laurent El Ghaoui. Randomized algorithms for the log-determinant optimization problem. 2015.

[91] Keith Kastella. Discrimination gain to optimize classification. *IEEE Transactions on Systems, Man, and Cybernetics Part A*, 1997.

[92] J. Frédéric Bonnans and Alexander Shapiro. *Perturbation Analysis of Optimization Problems*. Springer, 2000.

[93] Tyler W. Hughes, Momchil Minkov, Ian A. D. Williamson, and Shanhui Fan. Adjoint method and inverse design for nonlinear nanophotonic devices. *ACS Photonics*, 2018.

[94] Momchil Minkov, Ian A. D. Williamson, Lucio Claudio Andreani, Dario Gerace, Beicheng Lou, Alex Y. Song, Tyler W. Hughes, and Shanhui Fan. Inverse design of photonic crystals through automatic differentiation. *ACS Photonics*, 2020.

[95] Navin Khaneja, Timo Reiss, Cindie Kehlet, Thomas Schulte-Herbrüggen, and Steffen J. Glaser. Optimal control of coupled spin dynamics: design of NMR pulse sequences by gradient ascent algorithms. *Journal of Magnetic Resonance*, 2005.

[96] Ronald J. Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 8:229–256, 1992.

[97] John Schulman, Nicolas Heess, Theophane Weber, and Pieter Abbeel. Gradient estimation using stochastic computation graphs. In *NeurIPS*, 2015.

[98] Martin L. Puterman. *Markov Decision Processes: Discrete Stochastic Dynamic Programming*. Wiley, 1994.

[99] Dimitri P. Bertsekas and John N. Tsitsiklis. *Neuro-Dynamic Programming*. Athena Scientific, 1996.

[100] David H. Jacobson and David Q. Mayne. Differential dynamic programming. 1970.

[101] Yuval Tassa, Tom Erez, and Emanuel Todorov. Synthesis and stabilization of complex behaviors through online trajectory optimization. In *IROS*, 2012.

[102] David Silver and Joel Veness. Monte-Carlo planning in large POMDPs. In *NIPS*, 2010.

[103] Adrien Couëtoux, Jean-Baptiste Hoock, Nataliya Sokolovska, Olivier Teytaud, and Nicolas Bonnard. Continuous upper confidence trees. *LION*, 2011.

[104] Rémi Munos. From bandits to Monte-Carlo tree search: The optimistic principle applied to optimization and planning. *Foundations and Trends in Machine Learning*, 2014.

[105] Matthew Hausknecht and Peter Stone. Deep recurrent Q-learning for partially observable MDPs. In *AAAI Fall Symposium Series*, 2015.

[106] [CHECK] Zhong et al. [check] deep rl for sensor scheduling. 2018.

[107] C. C. Drovandi, J. M. McGree, and A. N. Pettitt. A sequential Monte Carlo algorithm to incorporate model uncertainty in Bayesian sequential design. *J. Comp. Graphical Stat.*, 23:3–24, 2014.

[108] Marcel Reginatto. Bayesian Cramér–Rao bound: A brief review. 2019.

[109] Fengwen Wang, Boyan S. Lazarov, and Ole Sigmund. On projection methods, convergence and robust formulations in topology optimization. *Structural and Multidisciplinary Optimization*, 2011.

[110] Ole Sigmund. Morphology-based black and white filters for topology optimization. *Structural and Multidisciplinary Optimization*, 2007.

[111] Boyan S. Lazarov, Fengwen Wang, and Ole Sigmund. Length scale and manufacturability in density-based topology optimization. *Archive of Applied Mechanics*, 2016.

[112] Rasmus E. Christiansen and Ole Sigmund. Compact 200 line MATLAB code for inverse design in photonics by topology optimization. *J. Opt. Soc. Am. B*, 2020.

[113] Shankar K. Selvaraja, Wim Bogaerts, Pieter Dumon, Dries Van Thourhout, and Roel Baets. Subnanometer linewidth uniformity in silicon nanophotonic waveguide devices using CMOS fabrication technology. *IEEE J. Sel. Top. Quantum Electron.*, 2010.

[114] Alexander Y. Piggott, Jesse Lu, Konstantinos G. Lagoudakis, Jan Petykiewicz, Thomas M. Babinec, and Jelena Vuckovic. Inverse design and demonstration of a compact and broadband on-chip wavelength demultiplexer. *Nature Photonics*, 2015.

[115] John M. Martinis and Anthony Megrant. *[CHECK] SQUID-based flux sensing.* 2014.

[116] Anthony J. Weiss and Ehud Weinstein. A lower bound on the mean-square error in random parameter estimation. *IEEE Trans. Inf. Theory*, 31, 1985.