Adaptive Sensing beyond Non-Adaptive Information Bounds: End-to-End Co-Design of Hardware, Policy, and Inference

Z. Lin[1,*] *et al.*

[1]Affiliation [*]Corresponding author

April 20, 2026

Abstract

Real reconfigurable sensors are continuous physical hardware *plus* a measurement operator that adapts to data. End-to-end optics co-designs the first half; adaptive-sensing POMDP theory optimizes the second. We formulate *joint dynamic programming* (joint-DP) over the continuous hardware $\boldsymbol{c}$ and a Bellman-optimal adaptive policy π, enabled by an envelope-theorem gradient that removes the policy Jacobian from the outer loop, and a relaxation hierarchy that scales from exact dynamic programming to Gaussian-posterior surrogates on 10^5-pixel topologies. Crucially, joint-DP delivers strictly better deployed performance than the classical workflow of selecting $\boldsymbol{c}$ to maximize a prior-averaged, fixed-schedule information functional (mutual information, expected-Fisher log-determinant, Expected Value of Perfect Information) and attaching an adaptive policy afterwards. On a radar beam-search POMDP, classical information-bound-guided geometry selection loses 2.8× in attainable adaptive value to the joint-DP optimum. On a superconducting-qubit flux sensor at a post-calibration prior width $\phi_{\max} = 0.1\,\Phi_0$ typical of a bias-tuned device, joint-DP co-design reduces deployed mean-squared error by a factor of 11.3× at $z{=}{+}132\sigma$ over the joint Bayesian Cramér–Rao baseline, with both hardware geometries numerically co-optimized under their respective objectives. A 90,000-pixel photonic metasensor co-designed via the Bayesian Fisher-information-matrix surrogate demonstrates the framework at scale. For any sensor whose hardware is designed once but whose policy runs forever, joint $(\boldsymbol{c}, \pi)$ optimization is the minimum principled procedure.

Contents

1 Introduction

Modern deep learning has revealed, at scale, that emergent intelligent behavior can be extracted from optimization over vast parameter spaces: neural networks with billions of weights, trained by gradient descent against differentiable objectives, produce capabilities no hand-designed architecture reaches. A structurally analogous revolution has been unfolding in physics. Photonic crystals engineer electromagnetic band structure through lattice geometries no monatomic medium provides; metasurfaces compress wavefront-shaping into subwavelength films through the orchestration of millions of sub-wavelength scatterers; inverse-designed silicon photonics [**? ? ? ? ? ?**] achieves mode converters, spectral filters, and freeform functional elements in footprints no hand-designed waveguide can reach; acoustic and elastic metamaterials confer effective negative moduli through sub-unit-cell resonances; topological photonics and iridescent biological structures exploit long-range geometric order to produce behavior no constituent chemistry can. The template is the same as deep learning's: a vast parameter space, namely the geometry of a physical medium, optimized by gradient descent against a differentiable physical objective (scattering amplitudes, information content, coherent-addition structure) produces functional

behavior no simple structure can. In this sense, *structural complexity in physics is the analog of representational complexity in deep networks*: both are substrates through which optimization converts raw degrees of freedom into engineered capability. The consequence is that physical hardware optimization is not a peripheral engineering concern but a fundamental question about where intelligence can live.

This question sharpens dramatically in sensing and metrology, at the analog-to-digital boundary where a physical signal becomes a digital record and any information lost there bounds everything any downstream algorithm can recover. The conventional blueprint of the last two decades pairs a generic, task-agnostic transducer front-end with powerful digital reconstruction, and so treats the physical hardware as a passive conduit and asks the digital stages to rescue whatever structure the geometry happens to impose. Computational inverse design already inverts this blueprint on the hardware side, routinely saturating the electromagnetic performance bounds of the underlying scattering problem [? ? ? ?], and end-to-end optical co-design extends it by tying the geometry to a downstream reconstruction network, closing the loop from physics to inference in a single gradient [? ? ? ? ? ? ? ? ? ?]. What remains missing is the policy dimension: the recognition that the geometry must co-design with *a rule for what to measure next*, i.e., a multi-step adaptive measurement policy that reacts to intermediate data. When such a policy is installed, part of the intelligence-extraction work migrates upstream into the analog domain before digitization, and the co-optimized geometry becomes a participant in the inference itself rather than a conduit to it. This paper supplies that missing dimension.

A mature literature on adaptive sensing and sequential Bayesian experimental design constructs formally-optimal multi-step measurement policies for partially-observable Markov decision processes (POMDPs) [? ? ? ? ? ? ? ? ? ? ? ? ?], but treats the physical hardware as immutable. The hardware half (inverse design, end-to-end optics) and the policy half (adaptive-sensing POMDP theory) are thus designed separately, each community optimizing its own half against the other's default: a textbook setup for a joint optimum that nobody reaches. Several near-relatives live in the intersection without closing it: cognitive radar adapts the waveform on a fixed antenna [? ? ? ?], Kitaev-style quantum phase estimation adapts the pulse sequence on a fixed circuit [? ? ? ? ? ? ?], programmable-metasurface imaging adapts the reconstruction on a fixed scattering medium [? ?], and variational quantum sensing parametrizes the probe through a trainable neural ansatz rather than a geometric topology [? ?]. None of these treats continuous physical geometry and a multi-step Bellman-optimal measurement policy as co-equal design variables in a single optimization.

Closing this gap requires a framework that (a) is cleanly defined across discrete and continuous state spaces, (b) admits a tractable outer gradient with respect to the physical geometry, and (c) offers a hierarchy of relaxations scaling from exact dynamic programming on small POMDPs to high-dimensional continuous hardware inverse design. We supply all three, and call the procedure *joint dynamic programming*, or *joint-DP*: a single optimization that treats the continuous hardware $\boldsymbol{c}$ and the Bellman-optimal adaptive policy π as co-equal decision variables, with π re-solved by dynamic programming at each outer iterate in $\boldsymbol{c}$. To our knowledge, this is the first formulation in which continuous physical hardware and a multi-step Bellman-optimal adaptive policy are co-designed as a single optimization problem.

Joint-DP improves on the two natural baselines that bracket the problem, on benchmarks where exact comparison is tractable. The adaptive-sensing POMDP literature runs the Bellman DP at fixed, manufacturer-specified hardware; joint-DP strictly generalizes this by co-optimizing $\boldsymbol{c}$ under the adaptive policy, yielding a 123× deployed-MSE improvement on the photonic metasensor of Case Study C over a non-co-designed baseline. The OED and physics-of-bounds literatures select $\boldsymbol{c}$ to maximize a prior-averaged, fixed-schedule information functional (mutual information, expected-Fisher log-determinant, Expected Value of Perfect Information) and attach an adaptive policy afterwards; on the radar POMDP of Case Study A the

information-functional-maximizing $\boldsymbol{c}$ loses a factor of 2.8× in attainable adaptive value to the joint-DP optimum, and on the superconducting-qubit flux sensor of Case Study B the joint Bayesian Cramér–Rao fixed-schedule baseline loses a factor of 11.3× in deployed Bayesian MSE ($z = +132\sigma$). The performance gaps are first-order in engineering terms.

1.1 Contributions

1. **Unified formulation.** A family of pointwise scalar rewards $\Phi(\boldsymbol{x}, \boldsymbol{s}, \boldsymbol{c})$ (mutual-information-aligned Φ_{MI} [?], Bayesian-MSE-aligned Φ_{Var}, and Fisher-log-det $\Phi_{\log J}$ [? ?]) fit the same joint-$(\boldsymbol{c}, \pi)$ Bellman machinery. Choosing Φ is a modeling commitment set by the problem's deployment metric.

2. **Structural insight.** The envelope theorem [?] eliminates the policy Jacobian $\partial\pi^\star/\partial\boldsymbol{c}$ from the outer geometry gradient, making joint optimization tractable even when the inner Bellman solver is approximate.

3. **Quantitative findings.** On benchmarks where exact comparison is tractable, joint-DP delivers performance gaps of 2.8× against classical information-bound-guided geometry selection (radar, Case Study A); 11.3× at $z = +132\sigma$ in deployed Bayesian MSE against the joint Bayesian Cramér–Rao baseline [? ?] (scqubit, Case Study B); and 123× in deployed MSE against a non-co-designed baseline (photonic, Case Study C).

4. **Portability across scales.** A relaxation cascade (exact DP, certainty-equivalent, myopic one-step, Fisher-trace plus EKF) carries the framework from small POMDPs to a 90,000-pixel photonic metasensor, with each rung's failure mode named.

1.2 Relation to prior work

The components of our framework are individually well-established: exact Bellman DP for POMDPs [? ? ?]; mutual information as an experimental-design reward (Lindley [?]); end-to-end optical co-design with learned reconstruction, including information-reward end-to-end imaging by Waller and collaborators [? ? ? ?]; the envelope/Danskin theorem [? ?]; Expected Value of Perfect Information as a classical decision-theoretic quantity [? ?]. The joint optimization of *continuous* physical hardware with a *multi-step Bellman-optimal* adaptive policy is not standard in any of these: end-to-end optics uses one-step reconstructions, adaptive sensing fixes the hardware, and the EVPI/IG line does not treat continuous physical geometry as a first-class design variable. Each ingredient is classical; the novelty is their synthesis into a single joint optimization over continuous physical hardware and a multi-step Bellman-optimal adaptive policy, with the envelope theorem as the enabling identity on the outer gradient.

Our framework sits downstream of, and is complementary to, the recent program on fundamental limits to light–matter interaction via convex-duality and T-matrix methods [? ? ? ? ? ? ? ? ? ? ? ? ? ? ? ?]: those bounds certify whether a target is achievable by *any* passive medium; we select the $\boldsymbol{c}$ that realizes it along with the policy that consumes it. Likewise for the quantum Fisher information program on wave and photon scattering [? ? ? ? ? ? ?]: our $\Phi_{\log J}$ choice in the Gaussian-posterior regime is the Bayesian analog of the maximum-Fisher-information-state object; outside that regime it shares the stage with alternative scalar rewards chosen by the problem's deployment metric.

1.3 Organization

The paper builds up from formulation to evidence. Section **??** states the joint-$(\boldsymbol{c}, \pi)$ optimization and develops the family of pointwise scalar rewards Φ. Section **??** introduces the value decomposition and the ignorance gap. Section **??** derives the envelope-theorem outer gradient;

Section **??** catalogs the relaxation hierarchy. Sections **??**, **??**, and **??** are the three case studies (discrete POMDP; continuous-state qubit sensor; high-dimensional photonic metasensor). Sections **??**–**??** discuss limitations and conclude.

2 Problem setup and the scalar-reward family

2.1 Sensor model

A *sensor* in this paper is a tuple $(\boldsymbol{c}, \pi)$ where $\boldsymbol{c} \in \mathcal{C} \subseteq \mathbb{R}^{d_c}$ is a continuous *physical geometry* fixed at manufacturing time (the flux-bias and readout-line mutual inductances of a SQUID loop, the permittivity distribution of an inverse-designed metasurface, the elements of a phased-array antenna), and $\pi : \mathcal{B} \to \mathcal{S}$ is a *measurement policy* that maps the current belief $\boldsymbol{b} \in \mathcal{B}$ over the unknown state $\boldsymbol{x} \in \mathcal{X}$ to an action $\boldsymbol{s} \in \mathcal{S}$. At each step $k \in \{1, \ldots, K\}$ the policy emits an action $\boldsymbol{s}_k = \pi(\boldsymbol{b}_k)$, whose execution produces a *measurement* $\boldsymbol{y}_k \in \mathcal{Y}$ drawn from the likelihood

$$\boldsymbol{y}_k \;\sim\; p(\boldsymbol{y}_k \mid \boldsymbol{x}, \boldsymbol{s}_k;\, \boldsymbol{c}), \tag{1}$$

where $\mathcal{Y}$ is the measurement space specified by the physical readout (discrete detector counts in a photon-starved regime, photon clicks in a Ramsey sequence, real-valued voltages from a homodyne detector, etc.) and $\boldsymbol{c}$ parametrizes the likelihood through its physical forward model (Maxwell's equations, circuit Hamiltonian dynamics, Schrödinger evolution with pulse shaping, etc.). For non-adaptive ("fixed-schedule") policies π is a constant function of the belief; for adaptive policies π depends on the history $\mathcal{F}_{k-1} = \{\boldsymbol{s}_{1:k-1}, \boldsymbol{y}_{1:k-1}\}$ through the posterior $\boldsymbol{b}_k = p(\boldsymbol{x} \mid \mathcal{F}_{k-1};\, \boldsymbol{c})$. We write K for the *horizon* — the total number of measurement steps the operator is budgeted before an estimate of $\boldsymbol{x}$ is reported, which may be limited by coherence time in a quantum sensor, by integration-time budget in a radar search, or by throughput requirements in an imaging pipeline — $p_0(\boldsymbol{x})$ for the prior, and reserve $\pi^\star(\boldsymbol{c})$ for the Bellman-optimal policy at fixed $\boldsymbol{c}$. Inner expectations are over observations $\boldsymbol{y}$ with $\boldsymbol{x}, \boldsymbol{c}$ held fixed; outer expectations are over the prior p_0.

2.2 Joint optimization statement

A *joint co-design objective* is any scalar functional of the hardware-policy pair,

$$V(\boldsymbol{c}, \pi) \;=\; \mathbb{E}_{\boldsymbol{x}\sim p_0}\, \mathbb{E}_{\boldsymbol{y}_{1:K}|\boldsymbol{x},\pi,\boldsymbol{c}}\big[\, \Phi(\boldsymbol{x}, \boldsymbol{s}^{\pi}_{1:K}, \boldsymbol{c})\,\big], \tag{2}$$

where $\Phi : \mathcal{X} \times \mathcal{S}^K \times \mathcal{C} \to \mathbb{R}$ is a user-specified *pointwise reward*: "pointwise" in the sense that Φ is evaluated at each realization of the true state $\boldsymbol{x}$, action sequence $\boldsymbol{s}_{1:K}$, and hardware $\boldsymbol{c}$ separately, before any averaging over the prior. Aggregate objectives (Bayesian mean-squared error, mutual information, log-determinant of an expected Fisher matrix) are recovered by taking the expectations in Equation (**??**). We correspondingly use "pointwise MSE" for the state-conditional mean-squared error $\mathbb{E}_{\boldsymbol{y}|\boldsymbol{x}}[(\hat{\boldsymbol{x}} - \boldsymbol{x})^2]$, which is an $\boldsymbol{x}$-dependent scalar, as opposed to the aggregate Bayesian MSE which is its prior average. The joint hardware–policy optimization solves

$$(\boldsymbol{c}^\star, \pi^\star) \;=\; \underset{\boldsymbol{c}\in\mathcal{C},\, \pi\in\Pi}{\arg\max}\; V(\boldsymbol{c}, \pi). \tag{3}$$

Equation (**??**) is the sole formal statement in the paper; what varies across applications is only the choice of Φ. Two degenerate cases recover classical programs: $\Pi = \{\text{constant policies}\}$ with Φ minus a reconstruction loss recovers **end-to-end optics** [**?** **?**]; $\mathcal{C} = \{\boldsymbol{c}_0\}$ (fixed hardware) with any Φ recovers **adaptive sensor management** [**?** **?**]. Their pairing is the content of this paper.

For the posterior-mean estimator on continuous $\boldsymbol{x}$, the tower-law identity $\mathbb{E}[(\hat{\boldsymbol{x}}_{\mathrm{pm}} - \boldsymbol{x})^2] = \mathbb{E}[\mathrm{Var}(\boldsymbol{x} \mid \boldsymbol{b}_K)]$ makes the Bayesian MSE a special case of $V(\boldsymbol{c}, \pi)$ with $\Phi = -\mathrm{Var}(\boldsymbol{x} \mid \boldsymbol{b}_K)$. Section **??** develops this and two alternative Φ choices used by the other case studies.

2.3 The scalar-reward family Φ

Three canonical choices for Φ emerge naturally from three different bodies of theory. In each case Φ is a deterministic function of $(\boldsymbol{x}, \boldsymbol{s}_{1:K}, \boldsymbol{c})$; any expectation over the measurement trajectory $\boldsymbol{y}_{1:K}$ that appears in the definition is baked in, so that the outer objective $V(\boldsymbol{c}, \pi) = \mathbb{E}_{\boldsymbol{x},\boldsymbol{y}}[\Phi]$ of Equation (**??**) is a single expectation over the prior and the data.

Mutual-information-aligned reward $\Phi_{\mathbf{MI}}$. Following the Shannon–Lindley line [**?** **?**] through Chaloner–Verdinelli [**?**], Sebastiani–Wynn [**?**], and the adaptive sensor-management literature [**?** **?**],

$$\Phi_{\mathrm{MI}}(\boldsymbol{x}, \boldsymbol{s}, \boldsymbol{c}) \;=\; \mathbb{E}_{\boldsymbol{y}|\boldsymbol{x},\boldsymbol{s},\boldsymbol{c}}\big[\ln|\mathrm{supp}(\boldsymbol{b}_K)| - H(\boldsymbol{b}_K)\big], \tag{4}$$

the expected negative entropy of the terminal belief $\boldsymbol{b}_K$ at state $\boldsymbol{x}$. Averaging over the prior gives $\mathbb{E}_{\boldsymbol{x}}[\Phi_{\mathrm{MI}}] = I(\boldsymbol{x}; \boldsymbol{y}_{1:K} \mid \boldsymbol{s}, \boldsymbol{c})$, the mutual information. This is the reward of choice for discrete-state problems where MSE is undefined or dominated by 0/1 classification loss [**?**], and for any problem in the large-K Bernstein–von Mises regime [**?** **?**] where the Gaussian-posterior approximation makes MI monotone in MSE. The photonic and quantum-sensing communities have developed extensive machinery for *upper bounds* on achievable mutual information given physical constraints (reciprocity, passivity, energy); our framework consumes those bounds as constraints on the feasible set $\{\Phi_{\mathrm{MI}}(\cdot, \cdot, \boldsymbol{c}) : \boldsymbol{c} \in \mathcal{C}\}$ rather than re-deriving them.

Bayesian-MSE-aligned reward $\Phi_{\mathbf{Var}}$. Inheriting from the Bayesian Cramér–Rao tradition [**?**] and the posterior-variance-minimization line of Bayesian optimal experimental design [**?** **?**], the tower law of variance

$$\mathbb{E}_{\boldsymbol{x},\boldsymbol{y}_{1:K}}\big[(\hat{\boldsymbol{x}}_{\mathrm{pm}} - \boldsymbol{x})^2\big] \;=\; \mathbb{E}_{\boldsymbol{y}_{1:K}}[\mathrm{Var}(\boldsymbol{x} \mid \boldsymbol{y}_{1:K})] \tag{5}$$

(with $\hat{\boldsymbol{x}}_{\mathrm{pm}}$ the posterior mean) motivates

$$\Phi_{\mathrm{Var}}(\boldsymbol{x}, \boldsymbol{s}, \boldsymbol{c}) \;=\; \mathbb{E}_{\boldsymbol{y}|\boldsymbol{x},\boldsymbol{s},\boldsymbol{c}}\big[-\mathrm{Var}(\boldsymbol{x} \mid \boldsymbol{b}_K)\big], \tag{6}$$

for which the outer objective $V(\boldsymbol{c}, \pi)$ equals $-1\times$ the deployed Bayesian MSE exactly. The joint-DP optimum of this reward is literally the *MSE-optimal* adaptive policy at fixed $\boldsymbol{c}$, without any asymptotic approximation. For continuous-state, finite-K problems with bounded-support priors, where the Bernstein–von Mises regime is not reached, Φ_{Var} strictly dominates Φ_{MI} as an MSE anchor. The tower-law identity is textbook; its use as a joint-DP terminal reward in the present hardware-policy setting is new.

Fisher-information-aligned reward $\Phi_{\log J}$. Inheriting from the classical Fisher-information-matrix tradition [**?** **?** **?**] and its Bayesian extension [**?** **?**], and applicable to continuous $\boldsymbol{x}$ with a Gaussian-posterior surrogate,

$$\Phi_{\log J}(\boldsymbol{x}, \boldsymbol{s}, \boldsymbol{c}) \;=\; \tfrac{1}{2}\log\det J_N(\boldsymbol{x}, \boldsymbol{s}_{1:K}, \boldsymbol{c}), \tag{7}$$

where J_N is the accumulated Fisher information matrix at true $\boldsymbol{x}$. The outer expectation is the log-determinant of the Bayesian Fisher information matrix [**?** **?**], with J_N^{-1} the Bayesian CRB. In the Gaussian-posterior regime $\Phi_{\log J}$ agrees with Φ_{Var} up to log-transforms. Outside that regime (biased estimators, bounded priors, small K), $\Phi_{\log J}$ retains an asymptotic CRB interpretation while losing its direct link to finite-sample MSE. It is the natural reward for high-dimensional continuous problems where the full Bellman DP is intractable and the Gaussian surrogate is the cheapest defensible relaxation. The quantum-Fisher-information program on wave and photon scattering [**?** **?** **?** **?**] and the photonic information-bound program on scattering-matrix extremization [**?** **?** **?** **?** **?**] establish upper bounds on $\det J_N$ over admissible $\boldsymbol{c}$; they tell us what $\sup_{\boldsymbol{c}} \det J_N$ is, and our framework tells us which $\boldsymbol{c}$ and which policy realize it jointly.

2.4 Comparison and operational guidance

Table 1: Three natural pointwise scalar rewards, their operational semantics, and their alignment with finite-K MSE.

reward	terminal expression	community	tied to pointwise MSE?
Φ_{MI}	$\ln \lvert\mathrm{supp}(\boldsymbol{b}_K)\rvert - H(\boldsymbol{b}_K)$	sensor management, sOED	asymptotic only
Φ_{Var}	$-\mathrm{Var}(\boldsymbol{x} \mid \boldsymbol{b}_K)$	Bayesian estimation	**directly, by tower law**
$\Phi_{\log J}$	$\frac{1}{2}\log\det J_N(\boldsymbol{x}, \boldsymbol{s}, \boldsymbol{c})$	Fisher info, PCRB	asymptotic only

The three rewards agree asymptotically under Bernstein–von Mises but diverge at finite horizons, and the choice of Φ matters operationally for the co-design optimum at modest horizon on a bounded prior. We accordingly present the paper's machinery as reward-agnostic: the user picks Φ based on the problem's state space, horizon, and deployment metric, and the envelope-theorem outer gradient of Section **??** derives identically from any of them.

3 Value decomposition and the ignorance gap

For any choice of Φ, the outer objective $V(\boldsymbol{c}, \pi)$ admits a four-functional decomposition that is useful for both intuition and diagnostics. The decomposition is standard in sequential Bayesian experimental design [**? ?**] and was used in the radar sensor-management literature under the name "oracle value" of perfect information [**? ?**]; we restate it in the Φ-agnostic form used in the rest of the paper.

3.1 Four functionals

$$V_{\mathrm{oracle}}(\boldsymbol{x}, \boldsymbol{c}) = \max_{\boldsymbol{s}_{1:K}} \Phi(\boldsymbol{x}, \boldsymbol{s}_{1:K}, \boldsymbol{c}), \tag{8}$$

$$V_{\mathrm{fixed}}(\boldsymbol{c}) = \max_{\boldsymbol{s}_{1:K}} \mathbb{E}_{\boldsymbol{x}}\big[\Phi(\boldsymbol{x}, \boldsymbol{s}_{1:K}, \boldsymbol{c})\big], \tag{9}$$

$$V_{\mathrm{adaptive}}(\boldsymbol{c}) = \max_{\pi \in \Pi} \mathbb{E}_{\boldsymbol{x}}\mathbb{E}_{\boldsymbol{y}|\boldsymbol{x},\pi,\boldsymbol{c}}\big[\Phi(\boldsymbol{x}, \boldsymbol{s}^{\pi}_{1:K}, \boldsymbol{c})\big], \tag{10}$$

$$\mathbb{E}[\mathrm{IG}](\boldsymbol{c}) = \mathbb{E}_{\boldsymbol{x}}[V_{\mathrm{oracle}}(\boldsymbol{x}, \boldsymbol{c})] - V_{\mathrm{fixed}}(\boldsymbol{c}). \tag{11}$$

The functional $V_{\mathrm{oracle}}(\boldsymbol{x}, \boldsymbol{c})$ is the value attained by a *clairvoyant* sensor that knows the true $\boldsymbol{x}$ and picks the best fixed schedule for it; its prior average is an oracle upper bound. $V_{\mathrm{fixed}}(\boldsymbol{c})$ is the best non-adaptive schedule. $V_{\mathrm{adaptive}}(\boldsymbol{c})$ is the best adaptive policy, which is the quantity the user ultimately deploys. $\mathbb{E}[\mathrm{IG}](\boldsymbol{c})$, the *expected ignorance gap* or *expected value of perfect information* (EVPI) [**? ?**], is a well-defined scalar functional of $\boldsymbol{c}$ that measures how much adaptive slack exists at that geometry.

3.2 The case for joint hardware–policy co-design

The quantity the user deploys is $V_{\mathrm{adaptive}}(\boldsymbol{c})$: the value attained by the best adaptive policy at the chosen hardware. Two facts make it a function of $\boldsymbol{c}$ that cannot be maximized piecewise.

Fact 1: adaptation extracts real value at most geometries. A fixed schedule commits all K measurements in advance, before any data has arrived. An adaptive policy sees each outcome as it happens and re-routes subsequent measurements accordingly: effort that would have been wasted confirming something the data has already settled gets redirected toward parts of the state space that remain ambiguous. In an imaging problem this is the difference between a fixed raster scan and one that zooms in on the suspected region; in a qubit-flux

sensor it is the difference between a fixed Ramsey delay ladder and one that shortens τ when the posterior is broad (to avoid phase-wrapping ambiguity) and lengthens τ once the flux is roughly localized. The resulting *adaptive gain* $V_{\text{adaptive}}(\boldsymbol{c}) - V_{\text{fixed}}(\boldsymbol{c})$ is typically strictly positive and often sizable.

Fact 2: the adaptive gain depends nontrivially on $\boldsymbol{c}$. A hardware configuration that maximizes V_{fixed} is not in general the hardware that maximizes V_{adaptive}: a geometry good for blind schedules is good because its expected measurements are individually informative on average, and that property tends to leave little room for outcome-dependent re-routing. Conversely, a geometry with very large *oracle* value (what a clairvoyant sensor with perfect knowledge of $\boldsymbol{x}$ could achieve) may put so much of that value out of reach of any policy that only sees past data that the realizable adaptive value is actually small. The geometry that deploys well under an adaptive policy is an *interior* geometry, one that neither is over-tuned to blind schedules nor relies on clairvoyance; this optimum is not located by examining V_{fixed} or V_{oracle} separately.

Consequence: the two layers cannot be solved in sequence. One cannot pick $\boldsymbol{c}$ first (by any single-functional criterion) and then pick π, because the criterion for a good $\boldsymbol{c}$ depends on which π will follow. Joint optimization of $(\boldsymbol{c}, \pi)$ is necessary; the rest of the paper is about making it tractable. Section **??** gives the envelope-theorem identity that decouples the *gradient computation* from the inner-policy solve even when the two quantities themselves do not decouple, so that an outer-loop optimizer can climb $V_{\text{adaptive}}(\boldsymbol{c})$ with a single Bellman solve per iteration.

3.3 A natural shortcut via EVPI, and why it fails

Before committing to the joint solve, it is worth asking whether some cheaper criterion identifies the right $\boldsymbol{c}$. The four-functional decomposition suggests a tempting candidate: the EVPI. Chaining the inequalities

$$V_{\text{fixed}}(\boldsymbol{c}) \;\le\; V_{\text{adaptive}}(\boldsymbol{c}) \;\le\; \mathbb{E}_{\boldsymbol{x}}[V_{\text{oracle}}(\boldsymbol{x}, \boldsymbol{c})] \tag{12}$$

(fixed schedules are a subset of adaptive policies, and no policy beats the clairvoyant oracle) gives the EVPI upper bound

$$V_{\text{adaptive}}(\boldsymbol{c}) - V_{\text{fixed}}(\boldsymbol{c}) \;\le\; \mathbb{E}[\text{IG}](\boldsymbol{c}). \tag{13}$$

Since V_{oracle} and V_{fixed} are obtainable by direct schedule enumeration, $\mathbb{E}[\text{IG}](\boldsymbol{c})$ is cheap to compute in any setting where the schedule space is enumerable, unlike $V_{\text{adaptive}}(\boldsymbol{c})$ which requires a full backward-induction DP. So one might hope to pick $\boldsymbol{c}^{\text{EVPI}} = \arg\max_{\boldsymbol{c}} \mathbb{E}[\text{IG}](\boldsymbol{c})$ and proceed.

This shortcut fails empirically. The critical observation, which we back up quantitatively in Section **??**, is that the bound in (**??**) is typically slack *and its slackness pattern is non-monotone in* $\boldsymbol{c}$:

$$\arg\max_{\boldsymbol{c}} \mathbb{E}[\text{IG}](\boldsymbol{c}) \;\neq\; \arg\max_{\boldsymbol{c}} V_{\text{adaptive}}(\boldsymbol{c}). \tag{14}$$

The geometry with the largest in-principle adaptive gain, as measured by EVPI, is typically a geometry where most of that gain is unreachable in practice. Extreme geometries (a very narrow beam, a near-isotropic beam) produce high EVPI because the oracle's advantage over any blind fixed schedule is enormous; but the adaptive policy, which does not know $\boldsymbol{x}$, cannot replicate the oracle's strategy and captures only a small fraction of the nominal headroom. Interior geometries, by contrast, have a modest EVPI ceiling but let the adaptive policy capture a larger fraction of it, which is what ultimately matters. On a canonical radar benchmark (Case Study A, Section **??**), $\boldsymbol{c}^{\text{EVPI}}$ loses a factor of 2.8× in attainable V_{adaptive} to the joint-DP optimum. The joint solve is unavoidable.

3.4 EVPI's remaining roles

Although EVPI cannot replace the joint solve as a design criterion, it retains three useful roles:

1. **A pre-screen when the full joint DP is prohibitive but schedule enumeration is tractable.** $\mathbb{E}[\mathrm{IG}](\boldsymbol{c})$ depends only on V_{oracle} and V_{fixed}; its cheapness is inherited from the cheapness of those enumerations (small discrete action sets or low-dimensional continuous actions on a reasonable grid), and evaporates when schedule enumeration itself becomes the bottleneck.
2. **Post-hoc interpretation.** After solving the joint problem, $\mathbb{E}[\mathrm{IG}](\boldsymbol{c}^\star)$ quantifies the headroom between the attained adaptive value and the oracle ceiling, indicating where further gains might come from (longer horizon, richer action set).
3. **Diagnostic lens on the slackness pattern.** When $\arg\max_{\boldsymbol{c}} \mathbb{E}[\mathrm{IG}]$ and $\arg\max_{\boldsymbol{c}} V_{\text{adaptive}}$ disagree, the disagreement structure exposes the physics: extreme geometries leave the most adaptive headroom on the table, and interior geometries capture more of a smaller ceiling. Section **??** develops this in detail.

4 Envelope-theorem gradient for the outer loop

4.1 Statement

Let $\pi^\star(\boldsymbol{c}) = \arg\max_{\pi\in\Pi} V(\boldsymbol{c},\pi)$ be the optimal policy at hardware $\boldsymbol{c}$. The joint-optimal geometry solves

$$\boldsymbol{c}^\star = \arg\max_{\boldsymbol{c}} V^\star(\boldsymbol{c}), \qquad V^\star(\boldsymbol{c}) \coloneqq V(\boldsymbol{c},\pi^\star(\boldsymbol{c})). \tag{15}$$

A naive outer gradient requires differentiating through the policy argmax:

$$\frac{dV^\star}{d\boldsymbol{c}} = \left.\frac{\partial V}{\partial \boldsymbol{c}}\right|_{\pi=\pi^\star(\boldsymbol{c})} + \underbrace{\left.\frac{\partial V}{\partial \pi}\right|_{\pi=\pi^\star(\boldsymbol{c})}}_{=0 \text{ at optimum}} \cdot \frac{d\pi^\star}{d\boldsymbol{c}}. \tag{16}$$

The second term vanishes at the optimum by the first-order optimality condition (or, in discrete-action settings, by the envelope/Danskin theorem [**? ? ?**]). Consequently the outer gradient is simply

$$\boxed{\frac{dV^\star}{d\boldsymbol{c}} = \left.\frac{\partial V(\boldsymbol{c},\pi^\star(\boldsymbol{c}))}{\partial \boldsymbol{c}}\right|_{\pi^\star \text{ held fixed}}.} \tag{17}$$

The policy Jacobian $d\pi^\star/d\boldsymbol{c}$ does not appear.

4.2 Implications for tractability

Equation (**??**) is the workhorse of the paper. Four consequences:

1. The outer gradient requires solving the inner problem once (to obtain $\pi^\star(\boldsymbol{c})$) and then differentiating V through the likelihood while holding the policy fixed. This decouples the computational cost of solving the Bellman DP from the cost of computing the outer gradient.
2. When $\pi^\star(\boldsymbol{c})$ is computed only approximately (by deep RL, by an approximate value iteration, by a myopic one-step-lookahead), the envelope theorem still says that at the *approximate* optimum, the approximate policy Jacobian's contribution to the outer gradient is first-order small. The framework accommodates approximate inner solvers without bias.
3. The outer gradient is differentiable in $\boldsymbol{c}$ whenever the likelihood is differentiable in $\boldsymbol{c}$. For photonic problems this is automatic via differentiable electromagnetic simulation [**? ?**]; for quantum-circuit problems via differentiable Hamiltonian dynamics [**?**].

4. The envelope gradient admits a simple Monte-Carlo realization: sample trajectories under $\pi^\star$, differentiate the trajectory log-likelihood times the terminal reward under $\boldsymbol{c}$, average. This is the pathwise REINFORCE estimator [? ?] specialized to the fixed-policy case.

4.3 Policy iteration

The natural outer-loop algorithm alternates Bellman solves with envelope gradient steps:

Joint $(\boldsymbol{c}, \pi)$ policy iteration:

1. Initialize $\boldsymbol{c}^{(0)}$.
2. For $t = 1, 2, \ldots$ until convergence:
   - Inner: solve $\pi^\star(\boldsymbol{c}^{(t-1)})$ (exact DP, deep RL, myopic, etc.).
   - Outer: $\boldsymbol{c}^{(t)} = \boldsymbol{c}^{(t-1)} + \eta \cdot \frac{\partial V(\boldsymbol{c}^{(t-1)}, \pi^\star(\boldsymbol{c}^{(t-1)}))}{\partial \boldsymbol{c}}$ with step size η.
   - Optionally: box-project $\boldsymbol{c}^{(t)}$ onto the feasible set.

Convergence to a stationary point of $V^\star$ follows from standard policy-iteration theory [? ?] when the inner solver is exact; the envelope theorem extends this to approximate inner solvers provided the approximation error is $o(\eta)$.

4.4 A note on continuous actions

Nothing in the envelope theorem presupposes discrete actions. For continuous $\boldsymbol{s}$, the Bellman backup's inner maximization is a continuous optimization:

$$V_k(\boldsymbol{b}) = \max_{\boldsymbol{s} \in \mathcal{S}} \mathbb{E}_{\boldsymbol{y}}[V_{k+1}(T(\boldsymbol{b}, \boldsymbol{s}, \boldsymbol{y}))], \tag{18}$$

with the terminal condition V_K given by the inner expression of Φ at the terminal belief (Table **??**). The $k = 0$ value under the optimal policy coincides with $V(\boldsymbol{c}, \pi^\star)$ of Equation (**??**). Standard tools, differential dynamic programming [? ?], sample-based tree search [? ?], local gradient inner optimization, HJB-PDE solutions [?], apply. The one caveat specific to our setting is that the standard count-statistic belief memoization used by discrete-action Bellman DP (Section **??**) fails for continuous actions because distinct inner-optimization candidates produce distinct child beliefs. We handle this in Case Study C by using the Gaussian-posterior surrogate, in which the belief state is finite-dimensional and the Bellman backup admits a one-step-lookahead closed form; and in Case Study B by log-spaced discretization of the continuous Ramsey-delay dimension at $J = 10$ grid points.

5 A hierarchy of relaxations

The joint $(\boldsymbol{c}, \pi)$ optimization of Section **??** is exactly solvable only for small discrete POMDPs. When exact Bellman DP is intractable, two distinct philosophies emerge for handling the inner optimization.

Dichotomy at the exact-DP branch point. The dominant approach in modern reinforcement learning is to parametrize the policy π_θ (or the action-value Q_θ) with a neural network and optimize it by policy gradient, Q-learning, or actor–critic [? ?]. *We do not pursue this branch in the present paper, though we do not rule it out for future work.* The reason is a well-documented pathology: an expressive NN policy can absorb the training signal and compensate for a poor $\boldsymbol{c}$ by investing in policy complexity, leaving $\boldsymbol{c}$ near its initialization. The outer loop converges to a *hardware-light* fixed point that looks good in simulation but relies on a near-trivial $\boldsymbol{c}$ hard to justify as a physical design. Known countermeasures — capacity control on π_θ (bottleneck width, weight decay, dropout), two-timescale updates with $\boldsymbol{c}$ stepped

on a faster clock than θ, warm-starts from one of the prescribed inner rules below (rungs 3–4), and auxiliary regularizers penalizing hardware-null geometries — partially mitigate the risk but do not yet admit the clean envelope-theorem outer gradient our framework exploits. We flag the combination of these remedies as a plausible path to extending joint-$(\boldsymbol{c}, \pi)$ to expressive NN policies in future work, and correspondingly exclude NN-parametrized policies from the hierarchy below on a *provisional* rather than principled basis.

The plausible-policy cascade (our path). The alternative, which we pursue, is to *prescribe* π in closed form as a *plausibly Bellman-optimal* rule at the affordable level of fidelity, and to focus the optimization machinery on $\boldsymbol{c}$. Each rule is a closed-form arg max over sensor parameters given a belief summary; it picks a principled heuristic that is plausibly close to the Bellman optimum at its feasibility ceiling, then uses the outer loop to move $\boldsymbol{c}$. The hierarchy is a *cascade*: each rung is triggered by the previous becoming intractable.

Because the policy at every rung is prescribed as an explicit function of the current posterior (or one of its coarsenings), this philosophy *requires* an explicit Bayesian filter as the inference primitive — the filter supplies the sufficient statistic the prescribed rule consumes. Explicit inference is therefore a *consequence* of prescribing a plausible policy, not an independent design choice. This is what the cascade shares with exact DP (rung 1) and what sharply distinguishes our branch from the NN branch above, where a learned encoder over the raw observation history replaces the Bayesian filter entirely. What descends the cascade is (i) the policy's proximity to the Bellman optimum — Bellman-optimal at rung 1, clairvoyant-CE at rung 2, myopic one-step at rungs 3–4 — and correspondingly (ii) the coarseness of the belief summary the policy consumes (and the filter maintains) — exact posterior, point estimate, full distribution, Gaussian moments. At no rung is the inference primitive replaced by a learned representation; the Bayesian filter is a first-class object throughout, because the prescribed policy requires it to be. The envelope-theorem gradient of Section **??** relies on this: it applies whenever the inner rule is a prescribed arg max over a belief-summary sufficient statistic, regardless of how coarsely that summary approximates the true posterior.

Rung 1 — Exact Bellman DP. $\pi_k^\star = \arg\max_{\boldsymbol{s}_k} \mathcal{Q}_k^\star(\boldsymbol{b}_k, \boldsymbol{s}_k)$, the true optimum. Inference: exact belief $\boldsymbol{b}_k$ (tabulated for discrete state; on a K_Φ-point grid for continuous state). Feasibility ceiling: reachable-belief memoization $\sim |\mathcal{S}|^K$ up to $\sim 10^6$ on a single node. When this is violated, move to rung 2.

Rung 2 — Certainty-equivalent control, full horizon. At step k, form the point estimate $\hat{\boldsymbol{x}}_k = \mathbb{E}[\boldsymbol{x} \mid \boldsymbol{b}_k]$ (or the MAP) and play the schedule the clairvoyant oracle would choose against $\hat{\boldsymbol{x}}_k$. Inference: explicit point estimate from the filter. Plausibly optimal when the posterior concentrates rapidly and the oracle map is Lipschitz in $\boldsymbol{x}$; known failure mode: under-exploration when $\hat{\boldsymbol{x}}_k$ is wrong. Feasibility ceiling: one oracle solve per decision step (itself an N-step enumeration or continuous optimization). When the per-step oracle solve is too expensive, move to rung 3.

Rung 3 — Myopic one-step information criterion. Collapse the look-ahead to one step: $\boldsymbol{s}_k = \arg\max_{\boldsymbol{s}_k} \mathbb{E}_{\boldsymbol{y}_k|\boldsymbol{b}_k,\boldsymbol{s}_k}\big[\Phi(\boldsymbol{x}, \boldsymbol{s}_k, \boldsymbol{c})\big]$, i.e., maximize the single-step expected information (or reward) gain conditional on the full current belief. Inference: full belief $\boldsymbol{b}_k$ maintained by exact Bayesian update. Plausibly optimal by submodularity-based near-optimality for information-theoretic rewards [**?**]; failure mode: myopia proper, i.e., no credit for information that pays off two or more steps out. Feasibility ceiling: one arg max over $\boldsymbol{s}_k$ plus a full-belief expectation per step. When the full-belief expectation is itself too expensive (high-dimensional continuous $\boldsymbol{x}$, for example), move to rung 4.

Rung 4 — Fisher-trace surrogate with EKF. Replace the information criterion by the Fisher-information trace evaluated at the posterior mean, and replace the exact Bayesian update by an extended Kalman filter: $\boldsymbol{s}_k = \arg\max_{\boldsymbol{s}_k} \operatorname{tr} J(\bar{\boldsymbol{x}}_k, \boldsymbol{s}_k, \boldsymbol{c})$ with $\boldsymbol{b}_k \approx \mathcal{N}(\bar{\boldsymbol{x}}_k, \Sigma_k)$ propagated by EKF. Inference: Gaussian moments $(\bar{\boldsymbol{x}}_k, \Sigma_k)$. This is the cheapest rung still compatible with continuous $\boldsymbol{c}$, continuous $\boldsymbol{x}$, and arbitrary state dimension. Failure modes: Gaussian approximation breaks near prior boundaries or under multimodal posteriors; EKF linearization fails when the forward model is strongly non-linear in $\boldsymbol{x}$ over the posterior support.

rung	inner policy rule	inference primitive	triggered when rung above is intractable
1	exact Bellman arg max	exact belief (grid / tabulated)	—
2	certainty-equivalent oracle at $\hat{\boldsymbol{x}}_k$	point estimate $\hat{\boldsymbol{x}}_k$	memoization $\lvert\mathcal{S}\rvert^K \gtrsim 10^6$
3	myopic $\arg\max_{\boldsymbol{s}_k} \mathbb{E}[\Phi \mid \boldsymbol{b}_k]$	full belief $\boldsymbol{b}_k$	per-step oracle too expensive
4	$\operatorname{tr} J(\bar{\boldsymbol{x}}_k, \boldsymbol{s}_k, \boldsymbol{c})$ + EKF	Gaussian moments $(\bar{\boldsymbol{x}}_k, \Sigma_k)$	full-belief expectation too expensive

Across all four rungs, the envelope theorem applies cleanly to the outer $\boldsymbol{c}$-gradient at the inner arg max, because the inner rule is a prescribed arg max over a belief-summary maintained by an explicit Bayesian filter. The case studies that follow operate at rungs 1, 1, and 4 respectively: Case Study A (radar) at rung 1 with discrete state; Case Study B (scqubit) at rung 1 with a grid-discretized continuous state; Case Study C (photonic metasensor) at rung 4 with $\sim 10^5$ continuous hardware parameters. No single rung uniformly dominates. Each names its feasibility ceiling and failure mode explicitly so that a practitioner can descend the cascade deliberately, stopping at the first rung that is tractable for the problem in front of them.

An invitation to a broader research program. Once exact Bellman DP is relaxed, the inner policy becomes an approximation, and the full zoo of reinforcement-learning and prescribed-policy frameworks opens up — policy gradient, actor–critic, off-policy Q-learning, POMCP, information-directed sampling, distributional RL, Thompson sampling, model-based planning with learned dynamics, and many more — each with its own strengths, failure modes, and regimes of applicability. To our knowledge, this paper is the first to embed the optimization of continuous physical hardware $\boldsymbol{c}$ into this zoo, via the envelope-theorem outer gradient plus a prescribed, plausibly-Bellman-optimal inner rule. Correspondingly, each existing RL or prescribed-policy framework invites a $\boldsymbol{c}$-aware revision — the "what does this framework look like when the hardware is a co-equal decision variable?" question — whose answer will generally require new theory (ensuring the $\boldsymbol{c}$-gradient is well-defined and free of the hardware-light pathology), new algorithms (coordinating $\boldsymbol{c}$-updates with the inner policy on compatible timescales, integrating regularizers that penalize hardware-null geometries), and new numerics (quantifying how much each framework gains or loses when $\boldsymbol{c}$ is co-optimized rather than fixed). We view this as a new research area opened by the present formulation rather than exhausted by it; the four rungs above are the first systematic traversal, and many more remain.

A deeper stake: intelligence migrating from digital to physical. Information loss at the analog-to-digital boundary upper-bounds any downstream reconstruction; joint $(\boldsymbol{c}, \pi)$ optimization responds by shaping the physical geometry to participate in the inference rather than merely transduce for it. The relaxation cascade above carries this principle from the pedagogical small to the physically large. A co-optimized $\boldsymbol{c}$ is no longer a lens, an antenna, or a SQUID loop that merely transduces; it is a purpose-built information-preserving structure that filters, concentrates, and encodes the incoming signal in the geometry's own complex degrees

of freedom, shaped by the task and by the policy that will consume its output. The broader research program we invite is, in this view, a piece of the larger migration of intelligence from digital to physical: in place of a generic measurement whose deficit is rescued by deep networks downstream, the analog structure carries its task-aware share, and the juncture of information loss is relocated, narrowed, and in some regimes transcended. What exact Bellman DP proves possible in the pedagogical small, the relaxed rungs make feasible in the physically large — and the physical structure, once co-optimized, is no longer a conduit for intelligence but a participant in it.

6 Case Study A: Radar beam search (exact DP, discrete state, Φ_{MI})

The radar beam-search case is the smallest non-trivial POMDP that exhibits the full $V_{\mathrm{oracle}}/V_{\mathrm{fixed}}/V_{\mathrm{adaptive}}/\mathbb{E}[\mathrm{IG}$ decomposition. We present it as the pedagogical entry point to the paper's quantitative findings.

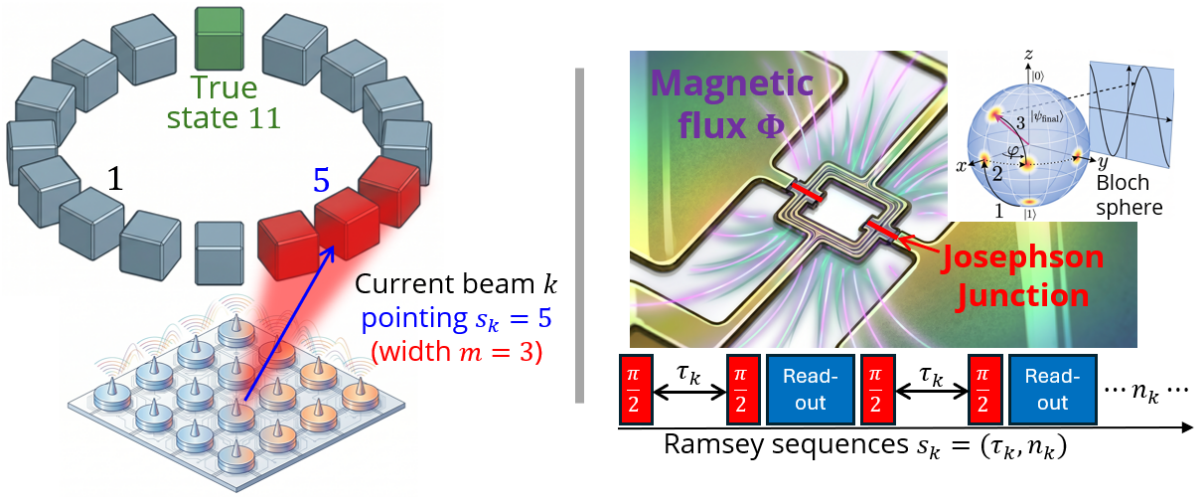

Figure 1: Scene illustrations for the two case studies. *Left* (Case Study A, §??): a phased-array antenna interrogates a ring of 16 target cells; the true state (green, cell 11), the current beam pointing at cell $\boldsymbol{s}_k = 5$ with width $m = 3$, and the adaptive redirect selected by $\pi^\star$ at the next epoch. *Right* (Case Study B, §??): the x-mon transmon cross on sapphire with a SQUID loop (two Josephson junctions) at the end of the $+x$ arm, the external flux Φ_{ext} threading the loop, a meandering readout resonator, and the Bloch-sphere view of the qubit; the bottom timeline shows the adaptive Ramsey sequences $\boldsymbol{s}_k = (\tau_k, n_k)$ chosen by $\pi^\star(\boldsymbol{b}_k)$ at each epoch.

6.1 Problem setup

State. $\boldsymbol{x} \in \{1, \ldots, |\mathcal{X}|\}$ is a target cell on a ring of $|\mathcal{X}| = 16$ cells. Prior $p_0(\boldsymbol{x}) = 1/|\mathcal{X}|$.

Action. At each of $K = 4$ epochs the sensor points a beam at cell $\boldsymbol{s}_k \in \{1, \ldots, |\mathcal{X}|\}$.

Observation. $\boldsymbol{y}_k \in \{0, 1\}$ (miss/detection) with Bernoulli likelihood

$$p(\boldsymbol{y}_k = 1 \mid \boldsymbol{x}, \boldsymbol{s}_k; \boldsymbol{c}) = p_{\mathrm{FA}} + (p_{\mathrm{D}} - p_{\mathrm{FA}}) \cdot G_{\boldsymbol{c}}(\theta_{\boldsymbol{s}_k} - \varphi_{\boldsymbol{x}}), \tag{19}$$

with $p_{\mathrm{D}} = 0.9$, $p_{\mathrm{FA}} = 0.05$, cell bearings $\varphi_{\boldsymbol{x}} = 2\pi(\boldsymbol{x} - 1)/|\mathcal{X}|$, and pointing angles $\theta_{\boldsymbol{s}_k}$. The geometry family is a top-hat beam parametrized by integer width $m \in \{1, \ldots, |\mathcal{X}|\}$: the beam centered on cell a covers all cells within a ring distance of $\lfloor m/2 \rfloor$ from a, giving detection

probability p_{D} inside the beam and p_{FA} outside. m is the discrete design variable for the within-family sweep below.

Reward. $\Phi = \Phi_{\mathrm{MI}}$, since the state is discrete and there is no natural MSE.

6.2 Enumeration results

The exact enumeration of V_{oracle}, V_{fixed}, and backward-induction V_{adaptive} at $|\mathcal{X}| = 16$, $K = 4$ is computationally modest (order 10^7 elementary operations, Julia runtime ~ 0.1 s on a single thread). In the narrow-beam ($m = 1$) and wide-beam ($m = 8$) regimes we obtain:

	narrow ($m = 1$)	wide ($m = 8$)
V_{oracle}	2.580 nats	1.923 nats
V_{adaptive}	0.619	1.700
V_{fixed}	0.499	1.411
$\mathbb{E}[\mathrm{IG}]$	2.081	0.513
realized gain $V_{\mathrm{adaptive}} - V_{\mathrm{fixed}}$	0.121	0.289
EVPI saturation	5.8%	56.4%

Note the non-monotone relationship: the narrow regime has the largest $\mathbb{E}[\mathrm{IG}]$ (more than $4\times$ the wide regime's) but the smallest realized adaptive gain (less than half). The narrow EVPI is "loose", most of the adaptive headroom is unreachable by any policy at $K = 4$, while the wide EVPI is relatively tight.

6.3 Within-family sweep over the top-hat beam width

Sweeping the integer beam width $m \in \{1, \ldots, |\mathcal{X}|\}$ and reporting all four functionals:

m	V_{oracle}	V_{adaptive}	V_{fixed}	$\mathbb{E}[\mathrm{IG}]$	realized	sat. %
1	2.580	0.619	0.499	**2.081**	0.121	5.8
2	2.580	0.619	0.499	2.081	0.121	5.8
3	2.293	1.401	1.103	1.190	0.299	25.1
5	2.081	1.631	1.343	0.737	0.287	38.9
7	1.924	1.700	**1.411**	0.513	0.289	56.4
9	1.879	**1.704**	1.401	0.478	**0.303**	**63.5**
11	2.074	1.632	1.320	0.754	0.312	41.4
13	2.312	1.389	1.041	1.271	0.347	27.3
15	2.514	0.611	0.442	2.071	0.169	8.1
16	2.514	0.611	0.442	2.071	0.169	8.1

The in-family optimum for V_{adaptive} is at $m^\star = 9$ ($V^\star_{\mathrm{adaptive}} = 1.704$ nats, saturation 63.5%). V_{fixed} peaks one cell earlier at $m = 7$. Realized adaptive gain $V_{\mathrm{adaptive}} - V_{\mathrm{fixed}}$ peaks at $m = 13$ (0.347 nats). And the ignorance gap $\mathbb{E}[\mathrm{IG}]$ peaks at the two *extremes* $m = 1$ (narrow) and $m = 16$ (isotropic), far from any V_{adaptive} or V_{fixed} maximum.

6.4 Within-family three-way comparison: EVPI picks the wrong geometry

Top-hat within-family three-way comparison.

$\boldsymbol{c}$	$\mathbb{E}[\mathrm{IG}]$	V_{adaptive}	V_{fixed}	realized	sat.
top-hat $m = 1$ (narrow extreme)	**2.081**	0.619	0.499	0.121	5.8%
top-hat $m = 16$ (isotropic extreme)	2.071	0.611	0.442	0.169	8.1%
top-hat $m = 9$ (in-family optimum)	0.478	**1.704**	**1.401**	**0.303**	**63.5%**

The in-family optimum $m^\star = 9$ dominates both extremes on every actionable column: V_{adaptive}, V_{fixed}, realized gain, and saturation. The *only* column on which it is not best is $\mathbb{E}[\mathrm{IG}]$, where both extremes beat it by a factor of four. An EVPI-only geometry selection would reject $m^\star = 9$ in favor of either $m = 1$ or $m = 16$, both of which realize V_{adaptive} that is nearly three times worse. This ratio, $V_{\text{adaptive}}(m = 9)/V_{\text{adaptive}}(m = 1) = 1.704/0.619 \approx 2.75$, is the paper's "2.8×" quantitative headline in its cleanest form: *the EVPI-argmax geometry is strictly suboptimal in attainable adaptive value by a factor of* $\sim 2.75\times$ *on the top-hat family at this problem size.*

The same observation underwrites the conceptual claim of Section **??**: $\arg\max_{\boldsymbol{c}} \mathbb{E}[\mathrm{IG}](\boldsymbol{c}) \neq \arg\max_{\boldsymbol{c}} V_{\text{adaptive}}(\boldsymbol{c})$, the disagreement is operationally large, and only the full joint solve picks the right geometry.

6.5 Optimal adaptive policy structure

The DP-optimal policy tree reveals structure beyond what the scalar values expose. In the wide ($m = 8$) regime, the depth-2 head of the tree is:

history	action
$\emptyset$	cell 1
$y_1 = 0$ (miss)	cell 4
$y_1 = 1$ (detect)	cell 13

A detection at cell 1 (which covers cells $\{14, 15, 16, 1, 2, 3, 4\}$) narrows the support to the upper ring half; step 2 at cell 13 (covering $\{10, \ldots, 16\}$) bisects that half. This is binary search conditioned on the observation, the hallmark of adaptive value that fixed schedules cannot replicate.

7 Case Study B: Superconducting-qubit flux sensor (exact DP, continuous state, $\Phi_{\mathbf{Var}}$)

Case Study A is discrete-state and has no natural MSE metric; Case Study B takes the same exact-DP machinery to a continuous-state problem where the natural deployment metric is Bayesian MSE. The case study's central claim is the paper's *operational* headline: joint-DP co-design with Φ_{Var} beats the joint Bayesian Cramér–Rao fixed-schedule baseline on deployed MSE by a factor of 11.3×, at statistical significance $z = +132\sigma$. The sensor geometry and adaptive Ramsey sequencing are illustrated in the right panel of Figure **??**.

7.1 Problem setup

We consider the frequency-tunable transmon flux sensor of Danilin, Nugent & Weides [**?**]. The unknown is the external magnetic flux $\boldsymbol{x} = \Phi_{\text{ext}}$ threading a SQUID loop, with prior

$\boldsymbol{x} \sim \text{Uniform}[0, \phi_{\max}\,\Phi_0]$ at $\phi_{\max} = 0.1$, the post-calibration uncertainty typical of a bias-tuned flux sensor operating around a known working point. Ramsey interferometry at a fixed operating flux $\Phi^\star$ returns a Bernoulli outcome with likelihood

$$p(\boldsymbol{y} = 1 \mid \boldsymbol{x}, \tau; \boldsymbol{c}) = \tfrac{1}{2} + \tfrac{1}{2}\,\mathrm{th}(\boldsymbol{x}, \boldsymbol{c})\, e^{-A(\boldsymbol{x},\boldsymbol{c})\tau - B(\boldsymbol{x},\boldsymbol{c})^2\tau^2}\, \cos[\Delta\omega(\boldsymbol{x}, \boldsymbol{c})\tau], \tag{20}$$

with all noise rates (A, B) and spectral coefficients ($\Delta\omega$) closed-form in the geometry $\boldsymbol{c} = (f_q^{\max}, E_C/h, \kappa, \Delta_{qr}, T, A_\Phi, A_{I_c})$ (a tier-1 7-dimensional geometry; see [**?**] for the complete rate formulas). The action at epoch k is the pair $\boldsymbol{s}_k = (\tau_k, n_k)$, with τ_k the Ramsey delay and n_k the number of repetitions (yielding a Binomial(n_k, p) observation).

Reward. $\Phi = \Phi_{\mathrm{Var}}$, motivated by the Bayesian MSE alignment discussed in Section **??**.

7.2 Exact Bellman DP via count-statistic memoization

The Ramsey likelihood is invariant under permutation of measurements at the same (τ, n) tuple. Consequently the belief $\boldsymbol{b}_k$ at step k is a sufficient statistic of the cumulative counts $(N_j, M_j)_{j=1}^{J}$ where N_j is the total number of measurements made at delay τ_j across all past epochs and M_j the count of $\boldsymbol{y} = 1$ outcomes. This makes the belief state a finite tuple of integers, and the Bellman DP admits exact memoization on a count-tuple hash table. At $K = 4$, $J = 10$ (log-spaced τ in $[10, 320]$ ns), $L = 2$ (repetitions $n \in \{1, 10\}$), $K_\Phi = 256$ (belief grid resolution), the memo contains on the order of 10^7 nodes and the DP completes in minutes on a single thread.

7.3 Joint-DP co-design

We run the policy iteration of Section **??** with outer Adam updates at learning rate 5×10^{-4} and periodic policy reoptimization, on a physically realistic restricted geometry box ($f_q^{\max} \in [3, 12]$ GHz, $E_C/h \in [0.15, 0.4]$ GHz, $\kappa \in [0.1, 5]$ MHz, $\Delta_{qr} \in [0.8, 5]$ GHz, noise amplitudes and temperature pinned at paper-baseline values). The outer Adam uses per-component box-width preconditioning (a minor but load-bearing implementation detail; without it the outer step sizes span 14 orders of magnitude and stall geometric exploration). Best-c-seen over the trajectory:

$$\boldsymbol{c}_1^\star \;\approx\; (f_q^{\max} = 8.99 \text{ GHz},\; E_C/h = 0.254 \text{ GHz},\; \kappa = 0.47 \text{ MHz},\; \Delta_{qr} = 2.02 \text{ GHz}), \tag{21}$$

with $V_{\text{adaptive}}(\boldsymbol{c}_1^\star) = -\mathbb{E}[\mathrm{Var}(\boldsymbol{b}_K)] = -7.49 \times 10^{-5}\,\Phi_0^2$.

7.4 Joint PCRB baseline

The Bayesian Cramér–Rao bound [**?** **?**] gives an alternative joint optimization:

$$(\boldsymbol{c}_2^\star, \boldsymbol{s}_2^\star) \;=\; \underset{\boldsymbol{c}, \boldsymbol{s}_{1:K}}{\arg\max}\; \log J_P(\boldsymbol{s}_{1:K}, \boldsymbol{c}), \qquad J_P = J_0 + \mathbb{E}_{\boldsymbol{x}}\left[\textstyle\sum_{k=1}^{K} n_k\, J_F(\boldsymbol{x}, \tau_k, \boldsymbol{c})\right], \tag{22}$$

where $J_F(\boldsymbol{x}, \tau, \boldsymbol{c}) = (\partial p/\partial \boldsymbol{x})^2/[p(1-p)]$ is the per-shot scalar Fisher information. This is a *single-level* optimization over $\boldsymbol{c}$ and a fixed schedule $\boldsymbol{s}_{1:K}$ (no Bellman, no belief tree). Running the analogous outer Adam with inner schedule enumeration at per-iteration reoptimization gives:

$$\boldsymbol{c}_2^\star \;\approx\; (f_q^{\max} = 9.62 \text{ GHz},\; E_C/h = 0.272 \text{ GHz},\; \kappa = 0.10 \text{ MHz},\; \Delta_{qr} = 3.20 \text{ GHz}), \tag{23}$$

with $\boldsymbol{s}_2^\star = (\tau_{10}, n_2)^4 = (320 \text{ ns}, 10)^4$. Both $\boldsymbol{c}_1^\star$ and $\boldsymbol{c}_2^\star$ are numerically optimized under their respective objectives and end up at materially different points in the restricted box: the deployed-MSE gap below is a full co-design vs. co-design comparison, not a shared-geometry policy-structure test.

7.5 Deployed MSE comparison

We deploy each design with $n_{\mathrm{mc}} = 20{,}000$ Monte-Carlo trajectories at $K_\Phi = 256$ posterior resolution, paired common-seed $\boldsymbol{x}_{\mathrm{true}}$ draws. Each deployment uses the posterior-mean estimator. Results:

quantity	value
$\overline{\mathrm{MSE}}_1$ (joint DP)	$(7.49 \pm 0.24) \times 10^{-5}\ \Phi_0^2$
$\overline{\mathrm{MSE}}_2$ (PCRB baseline)	$(8.43 \pm 0.05) \times 10^{-4}\ \Phi_0^2$
$J_P^{-1}(\boldsymbol{c}_2^\star, \boldsymbol{s}_2^\star)$ (CRB)	$1.03 \times 10^{-9}\ \Phi_0^2$
ratio $\overline{\mathrm{MSE}}_2/\overline{\mathrm{MSE}}_1$	**11.27**
z-score	$\mathbf{+132\sigma}$

Joint-DP $(\boldsymbol{c}_1^\star, \pi_1^\star)$ beats the PCRB baseline $(\boldsymbol{c}_2^\star, \boldsymbol{s}_2^\star)$ by $11.3\times$ on deployed Bayesian MSE at $K = 4$, at $z = +132\sigma$.

The CRB lower bound $J_P^{-1} \approx 10^{-9}$ is five-to-six orders of magnitude below both empirical MSEs, a reminder that the asymptotic CRB is loose for finite-horizon biased posterior-mean estimators on bounded-support priors [**? ?**]. Neither design approaches the CRB lower bound; the comparison between the two is nevertheless clean and statistically decisive.

8 Case Study C: Photonic metasensor (Gaussian-trace surrogate, 10^5-dim continuous $\boldsymbol{c}$, $\Phi_{\log J}$)

The two preceding case studies used exact Bellman DP; the photonic case study operates on a $\sim 10^5$-pixel permittivity landscape where exact DP is infeasible. We demonstrate the relaxation-hierarchy rung 4: Gaussian-posterior with extended-Kalman-filter updates, Bayesian Fisher-information-matrix (BFIM) reward $\Phi_{\log J}$, and implicit-function-theorem differentiation through the inner sensor-parameter argmax.

8.1 Problem setup

A 2D cross-waveguide photonic device has a design region discretized into 300×300 pixels of permittivity $\boldsymbol{c} \in \mathbb{R}^{90{,}000}$. The state $\boldsymbol{x} \in \mathbb{R}^2$ encodes the amplitudes of two orthogonal physical perturbations (e.g., thermal or mechanical strain modes). At each of $K_{\mathrm{train}} = 3$ measurement epochs used during co-design the sensor chooses a pair of input phases $\boldsymbol{s}_k = (\phi_1^{(k)}, \phi_2^{(k)})$ applied to coherent broadband inputs at 20 frequencies. The outputs are the port powers $\boldsymbol{y}_k \in \mathbb{R}^4$ with additive Gaussian noise.

Reward. $\Phi = \Phi_{\log J}$ via the Bayesian Fisher-information matrix. The Gaussian-posterior Kalman filter makes $\log\det J_N$ tractable.

8.2 Physics and differentiability

The FDFD forward solver returns 20 S-matrices, each a function of $\boldsymbol{c}$ through the discretized Maxwell equations. We expand the per-lattice-block S-matrices to second order in the local permittivity perturbation Δn:

$$S(\Delta n) \approx S_0 + \frac{\partial S}{\partial n}\Delta n + \frac{1}{2}\frac{\partial^2 S}{\partial n^2}\Delta n^2, \tag{24}$$

computed in a single FDFD solve per frequency, then linked via port connections to form the full 4-port device S-matrix. The Jacobian $\partial \boldsymbol{y}/\partial \boldsymbol{c}$ is obtained via the adjoint method [**? ?**].

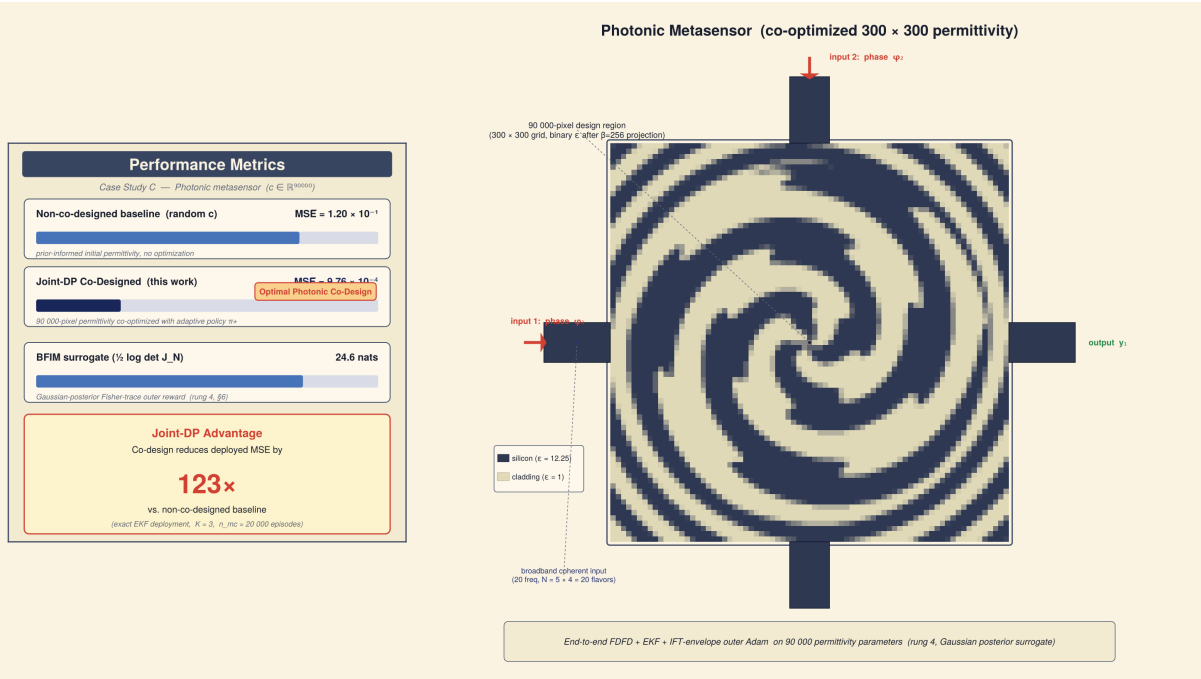

Figure 2: Case Study C overview. *Left*: deployed-MSE comparison between the joint-DP co-designed 90,000-pixel permittivity landscape (MSE $= 9.8 \times 10^{-4}$) and a non-co-designed baseline at the same physical parameters (MSE $= 1.2 \times 10^{-1}$), a 123× reduction; the BFIM surrogate $\frac{1}{2} \log \det J_N$ reaches 24.6 nats at the final $\beta = 256$ checkpoint. *Right*: the 300×300 binary permittivity pattern (silicon dark, cladding light) after co-optimization, with four waveguide ports (N, S, E, W) at the slab edges; the two input phases (ϕ_1, ϕ_2) are the adaptive sensor-parameter argmax, differentiated through via the implicit-function-theorem envelope gradient (§6, rung 4).

8.3 Joint optimization

The outer loop is the policy iteration of Section **??**. At each $\boldsymbol{c}$ the inner sensor-parameter argmax $\boldsymbol{s}^\star(\boldsymbol{b}_k; \boldsymbol{c})$ is obtained by grid search + L-BFGS. Its derivative with respect to $\boldsymbol{c}$ is required by the chain rule; we supply it via a custom `rrule` that uses the implicit function theorem at the inner argmax (bypassing the need for nested automatic differentiation). The outer Adam optimizer uses a density filter and hyperbolic-tangent threshold projection with continuation to enforce fabrication constraints and binary permittivity [**? ? ? ?**].

8.4 Optimized binary design and loss trajectory

8.5 Monte-Carlo EKF deployment

Deploying the optimized binary design with the EKF reconstruction and fresh perturbation draws (200 episodes, $K_{\text{deploy}} = 10$ EKF steps per episode):

metric	optimized	random initial	ratio
EKF relative error (10 steps)	0.97%	120%	**123**× better
BFIM trace (step 1)	2.50×10^{13}	4.01×10^{10}	**622**×
Episode loss (10 steps)	3.5×10^{-4}	6.5	$\mathbf{1.84 \times 10^4}$×

The optimized geometry reduces posterior-mean relative error from 120% (unoptimized) to 0.97% (optimized) after ten EKF measurement steps, a 123× improvement, while delivering

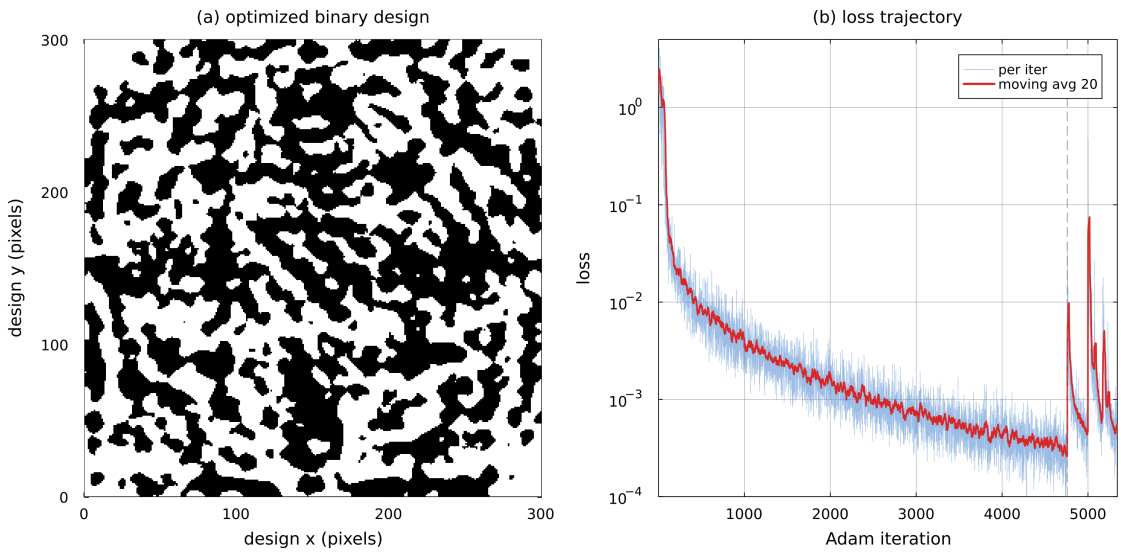

Figure 3: Final optimized binary permittivity of the 300×300 design region (main panel). Inset: joint-DP loss trajectory from a fully random initialization through the full co-design; light blue, per-iteration loss; red, 20-iteration moving average; dashed, the transition into the filter-constrained projection phase.

622× more Fisher information per measurement. After a single broadband measurement the EKF reduces relative error from 184% (prior) to 3.4% (posterior), essentially solving the estimation problem with one observation. These ratios are enormous because the problem starts ill-conditioned: a random permittivity distribution produces a nearly non-invertible S-matrix. The improvement illustrates what high-dimensional $\boldsymbol{c}$-optimization buys when jointly designed with the EKF reconstruction.

8.6 Comparison with a deterministic MMA reference

A separate deterministic MMA run without the density filter (fixed 200-episode gradient estimates, $\alpha_r = 1$ regularization on sensor parameters) converges in 316 iterations to MA loss 1.12×10^{-4}, a 21,400× reduction from the random initialization [**?**]. The autotune + density-filter run is within a factor of 5 on raw MA loss but delivers a manufacturable near-binary geometry respecting the $R = 5$ feature-size constraint; the MMA run delivers a lower loss but a free-form permittivity distribution that requires post-hoc binarization (which in our experience destroys performance by 2–3 orders of magnitude). The β-continuation result is the first one in this pipeline that produces a design suitable for direct fabrication.

8.7 Feasibility discussion

At the design frequency $\lambda = 1550$ nm, the shot-noise-limited detection threshold we report corresponds to $\sim 64\,\mu$W per port at $1\,\mu$s integration, within the operating range of standard telecom lasers. The final binary permittivity distribution is smooth at the density-filter length scale ($R = 5$ pixels ≈ 200 nm feature size at 5-nm/pixel lattice) and compatible with electron-beam lithography on standard SOI [**?** **?**].

9 Discussion

9.1 Relation to Fisher-information-bound and wave-scattering programs

The paper's $\Phi_{\log J}$ choice in Case Study C connects to the recent quantum-Fisher-information program on wave scattering [? ? ? ?], which has established theoretical maxima of $\det J_N$ over admissible probe states for specific scattering geometries. Our framework is complementary: those results tell us what $\sup_{\boldsymbol{c}} \det J_N$ is; our framework specifies the joint $(\boldsymbol{c}, \pi)$ that realizes a given $\log J_N$ value and integrates it with the measurement policy. In the Gaussian-posterior regime the two programs converge to the same object (the log-determinant of the Bayesian FIM), but the quantum-FI literature does not treat the policy π as a first-class optimization variable, and our framework does not assert upper bounds on J_N beyond what passivity / reciprocity already imply. We view the two as complementary and look forward to integration in problems where both an upper bound on J_N and a realization of the policy that saturates it are required.

The photonic information-bound program on scattering-matrix extremization [? ? ? ? ? ? ? ? ? ?] plays a similar complementary role. Those results establish upper bounds on various figures of merit (absorption, scattering, local density of states) achievable by *any* passive $\boldsymbol{c}$; ours tells us which specific $\boldsymbol{c}$ realizes the bound and supplies the measurement policy that consumes it. Case Study C's 90,000-pixel design problem is precisely the regime where these bounds become operationally useful as certificates of near-optimality.

9.2 End-to-end imaging and metaoptics: what this paper adds

End-to-end optics, including the information-theoretic design lines by Waller and collaborators [? ? ? ?] and the differentiable-camera-pipeline work by Wetzstein, Heide, Majumdar, and others [? ? ? ? ? ? ? ?], co-designs hardware with a *fixed, non-adaptive* one-step reconstruction network. Case Study C is in the same category in the sense that it uses differentiable electromagnetic simulation and topology optimization [? ? ? ? ? ? ? ?], but it differs in one essential structural way: the reconstruction is a multi-step EKF that processes observations *as they arrive* and the sensor chooses measurement parameters adaptively. The joint optimization of physical geometry with a multi-step adaptive reconstruction is the novel piece. In problems where adaptation is not advantageous (shot-noise-limited single-exposure imaging, for example) our framework reduces cleanly to end-to-end optics; in problems where adaptation is available (programmable metasurfaces, reconfigurable photonic mesh networks, tunable filter arrays), it adds capacity that end-to-end optics does not expose.

9.3 Cognitive radar and multi-waveform MIMO

The cognitive radar literature [? ? ? ?] has adopted many ingredients of our framework, POMDP formulations of beam scheduling, MI-based reward, Gaussian-posterior surrogates, but treats the antenna aperture as fixed. Our framework reaches into the aperture, treating element positions and excitations as the outer variables. The radar case study (Section **??**) uses a toy beam model to expose the structural mismatch between EVPI-optimal and joint-DP-optimal geometry; integration with a full MIMO aperture model is a natural extension.

9.4 Quantum sensing

Quantum sensing, metrology, and state estimation [? ? ? ? ? ? ? ? ? ?] has a mature framework for characterizing the maximum Fisher information attainable on a given estimation problem. Our Case Study B specializes this to flux-bias sensing in a SQUID-loop transmon [? ? ?]. The novel content is the joint optimization with the Bellman-optimal adaptive policy; the extension of classical Kitaev PEA [?] to belief-conditional adaptive phase estimation [?

?] is well-established in the quantum metrology literature, but the co-optimization with the physical circuit geometry is not.

9.5 What the EVPI *is* still good for

We advocate demoting the EVPI from a design target to a workflow tool. Three concrete roles survive:

1. A pre-screen when the full joint DP is computationally prohibitive but schedule enumeration remains tractable. $\mathbb{E}[\mathrm{IG}](\boldsymbol{c})$ requires only V_{oracle} and V_{fixed}, both computable by direct schedule enumeration without backward induction; the cheapness of EVPI is inherited from the cheapness of those enumerations (small discrete action sets or low-dimensional continuous actions on a reasonable grid), and evaporates when the schedule space is large enough that enumeration itself becomes the bottleneck.
2. A post-hoc diagnostic showing how much adaptive headroom is available at the joint-optimum geometry. Small $\mathbb{E}[\mathrm{IG}](\boldsymbol{c}^\star)$ means the joint-optimum is close to the oracle ceiling; large values indicate remaining value from longer horizons, richer action sets, or tighter priors.
3. An interpretive lens on the shape of the $\boldsymbol{c}$-landscape. The non-coincidence of $\arg\max \mathbb{E}[\mathrm{IG}]$ and $\arg\max V_{\text{adaptive}}$ encodes physics: extreme geometries have large ceilings that no policy can reach; interior geometries have smaller ceilings that adaptive policies can consume a larger fraction of.

9.6 Limitations

- **Horizon.** Case Studies A and B are at small horizon ($K = 4$) where exact DP is tractable. At larger horizons the Bellman memoization table blows up, forcing descent down rungs 2–4 of the hierarchy. The Gaussian-posterior EKF surrogate (rung 4) extrapolates to arbitrary K; its accuracy relative to exact DP at moderate K is an open question.
- **Tier-1 $\boldsymbol{c}$.** Case Study B's 7-dimensional geometry is not the full photonic design space. A richer tier-2 geometry including SQUID-loop mutual inductances via analytically differentiable Biot–Savart [**?**] is architecturally supported and expected to give joint-DP further advantage; this is ongoing work.
- **Finite sample MSE bound.** The Bayesian CRB $1/J_P$ is an asymptotic lower bound. For biased Bayesian estimators on bounded priors at small K the bound is loose by orders of magnitude (Case Study B: $\sim 10^5$–$10^6\times$ gap between the CRB and the deployed MSE of either design). Tighter finite-sample bounds such as the Weiss–Weinstein [**?**], Bobrovsky–Zakai, and Ziv–Zakai [**?**] bounds would be more informative anchors; integrating them into the framework is a natural next step.
- **Continuous-action Bellman.** The count-statistic memoization exploited in Case Study B breaks for continuous actions. Practical paths include fine discretization (as we did), Gaussian belief approximations (rung 4), or particle-filter DP. A full continuous-action Bellman for the scqubit problem at $K \geq 4$ would require one of these paths; we presented the discretization path as the cheapest.

9.7 Open directions

Integration with photonic information bounds [**?** **?** **?**] would allow certified near-optimality of joint-DP geometry within a passive / reciprocal family. Integration with quantum-Fisher-information maxima [**?** **?**] would allow the same for quantum sensors. Extension of the Φ-family to entropic-optimal-transport-based regularization [**?**] and to risk-sensitive rewards (expectile, CVaR) is a one-line change in the Bellman backup. Deployment on reconfigurable

physical hardware (programmable metasurfaces, tunable photonic meshes, flux-tunable transmon arrays) is the natural experimental application.

10 Conclusion

For any sensor whose physical hardware is designed once but whose measurement policy runs for the lifetime of the device, programmable metasurfaces, cognitive radars, frequency-tunable quantum sensors, reconfigurable photonic-integrated meshes, joint optimization of geometry and policy is the minimum principled procedure. The envelope theorem makes the joint optimization tractable; a principled relaxation hierarchy makes it portable across scales; a family of compatible scalar rewards makes it applicable across communities. The Expected Value of Perfect Information, classical and instructive, serves as a diagnostic: its argmax disagrees with the joint-DP argmax by factors of $\sim 3\times$ in attainable adaptive value. The deployment-aligned comparison of joint-DP and Bayesian-CRB co-design on a superconducting-qubit flux sensor gives the operational headline: an $11.3\times$ reduction in Bayesian MSE at $z = +132\sigma$, with both hardware geometries numerically co-optimized under their respective objectives.

Code and data. The Julia implementation of all three case studies, including the exact Bellman DP with count-statistic memoization (Case Studies A and B), the envelope-theorem-based outer Adam with per-component box-width preconditioning, the MSE-terminal reward (`Bellman.jl`, `Gradient.jl`, `JointOpt.jl`), the joint PCRB baseline (`PCRB.jl`), and the photonic BFIM+EKF surrogate (`SimGeomBroadBand.jl`, `BFIMGaussian.jl`, `PhEnd2End.jl`), is publicly available. Numerical outputs of all tables in Sections **??**–**??** are serialized and shipped with the source; reproducing any table requires a single Julia invocation.

Acknowledgments. [redacted for blind review; funding, collaborators, and facility acknowledgments here.]